

# Recommender Systems

## Lecture 4

# Previous lecture

## Lifecycle of a recsys experiment

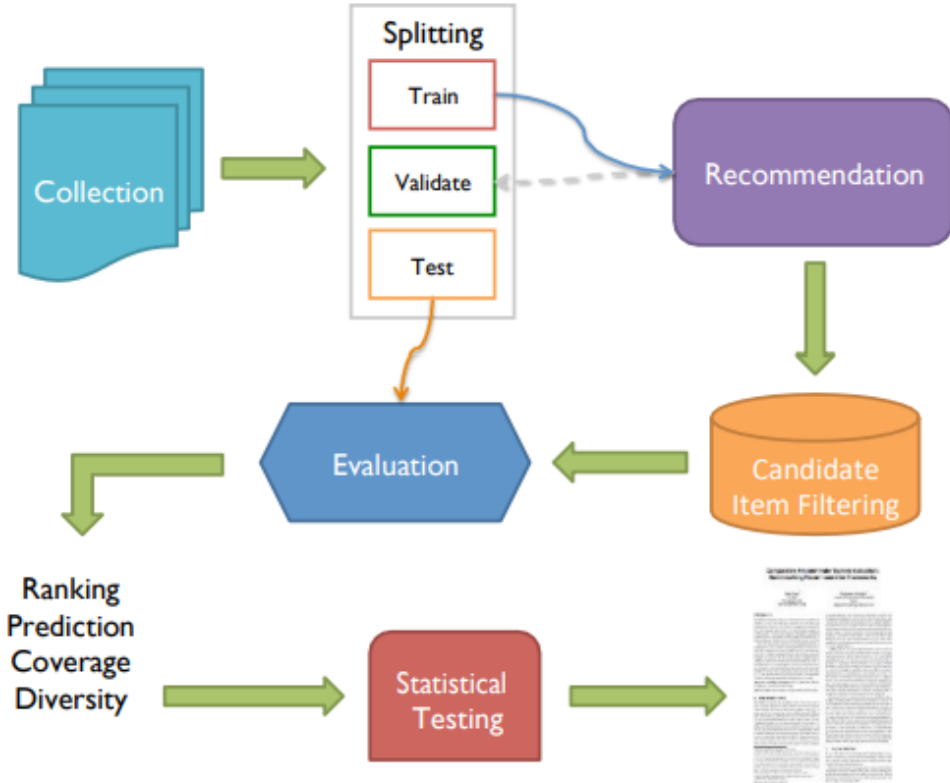



Image source: Bellogín, Alejandro, and Alan Said. "Improving accountability in recommender systems research through reproducibility." User Modeling and User-Adapted Interaction (2021): 1-37.

## Data splitting



### Splitting options:

- entry- or user-wise
- warm start
  - strong/weak generalization

## Holdout sampling

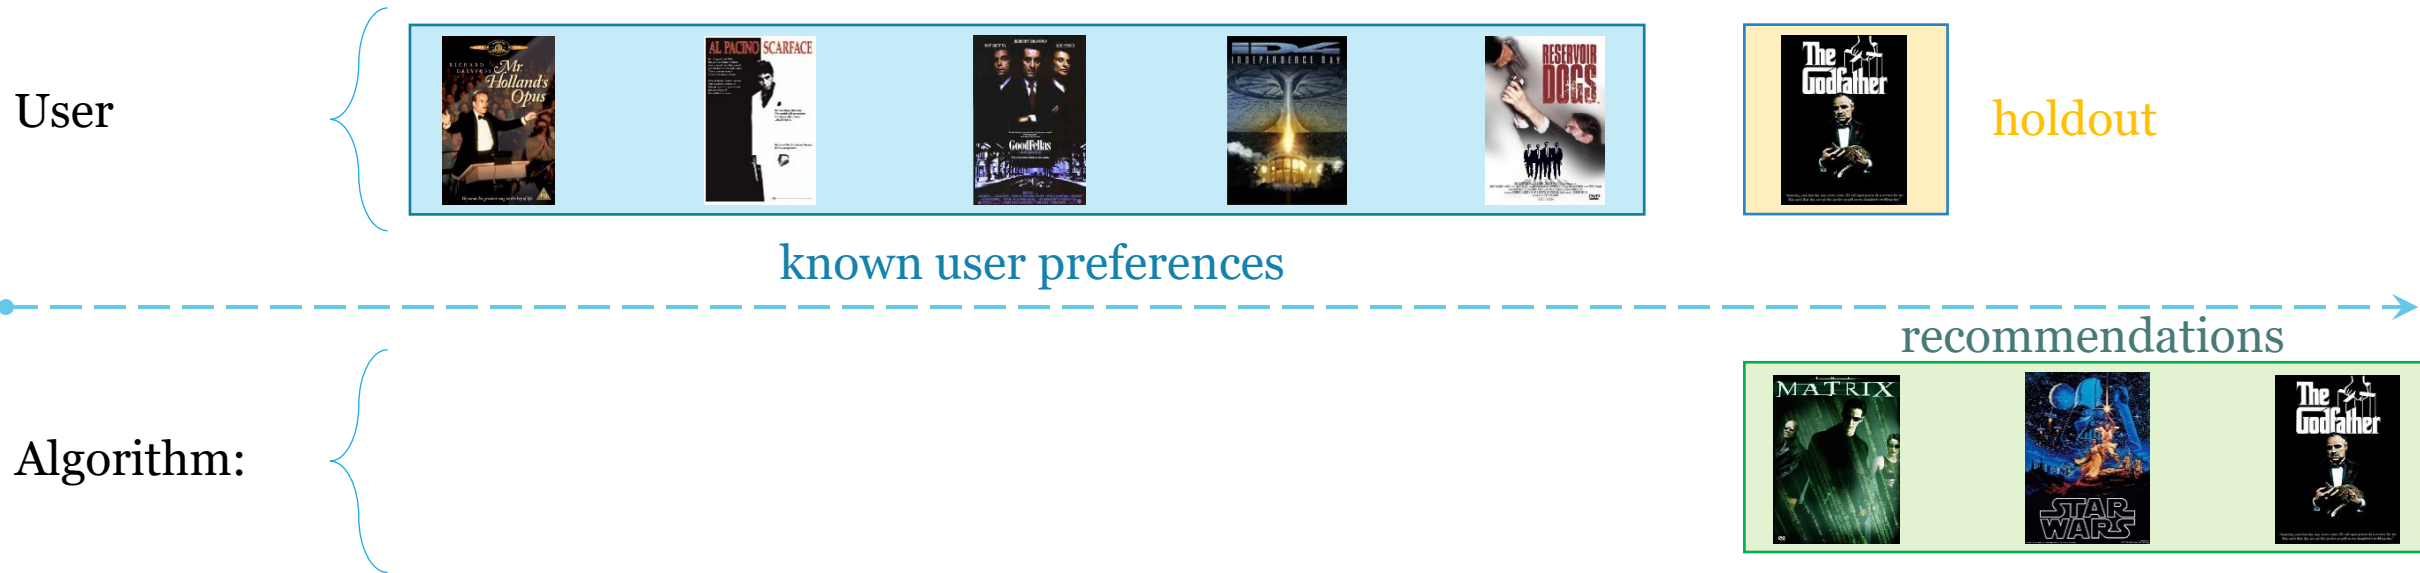
### Strategies

- Random
- Rating-based (e.g., top-rated)
- Temporal (e.g., most recent)

### Sample size

- Fixed number of items (e.g., 1)
- Fixed percentage of items

# Previous lecture



$$HR = \frac{1}{\#(\text{test users})} \sum_{\text{test users}} \text{hit}$$

$$\text{hit} = \begin{cases} 1 & \text{if holdout item is in recommended items,} \\ 0 & \text{otherwise.} \end{cases}$$

$$MRR = \frac{1}{\#(\text{test users})} \sum_{\text{test users}} \frac{1}{\text{hit rank}}$$

hit rank = position of the item in the recommendations list

Typically computed: metric@n, where  $n$  = #recommended items, e.g. Recall@n, MRR@n, etc.

# Task

Consider top-2 recommender for 10 users from 20-items (2 items per user).

**What's better:**

1) Correctly recommend 2 items to each of 5 users?

or

2) 1 item to each of 10 users?

**Why?**

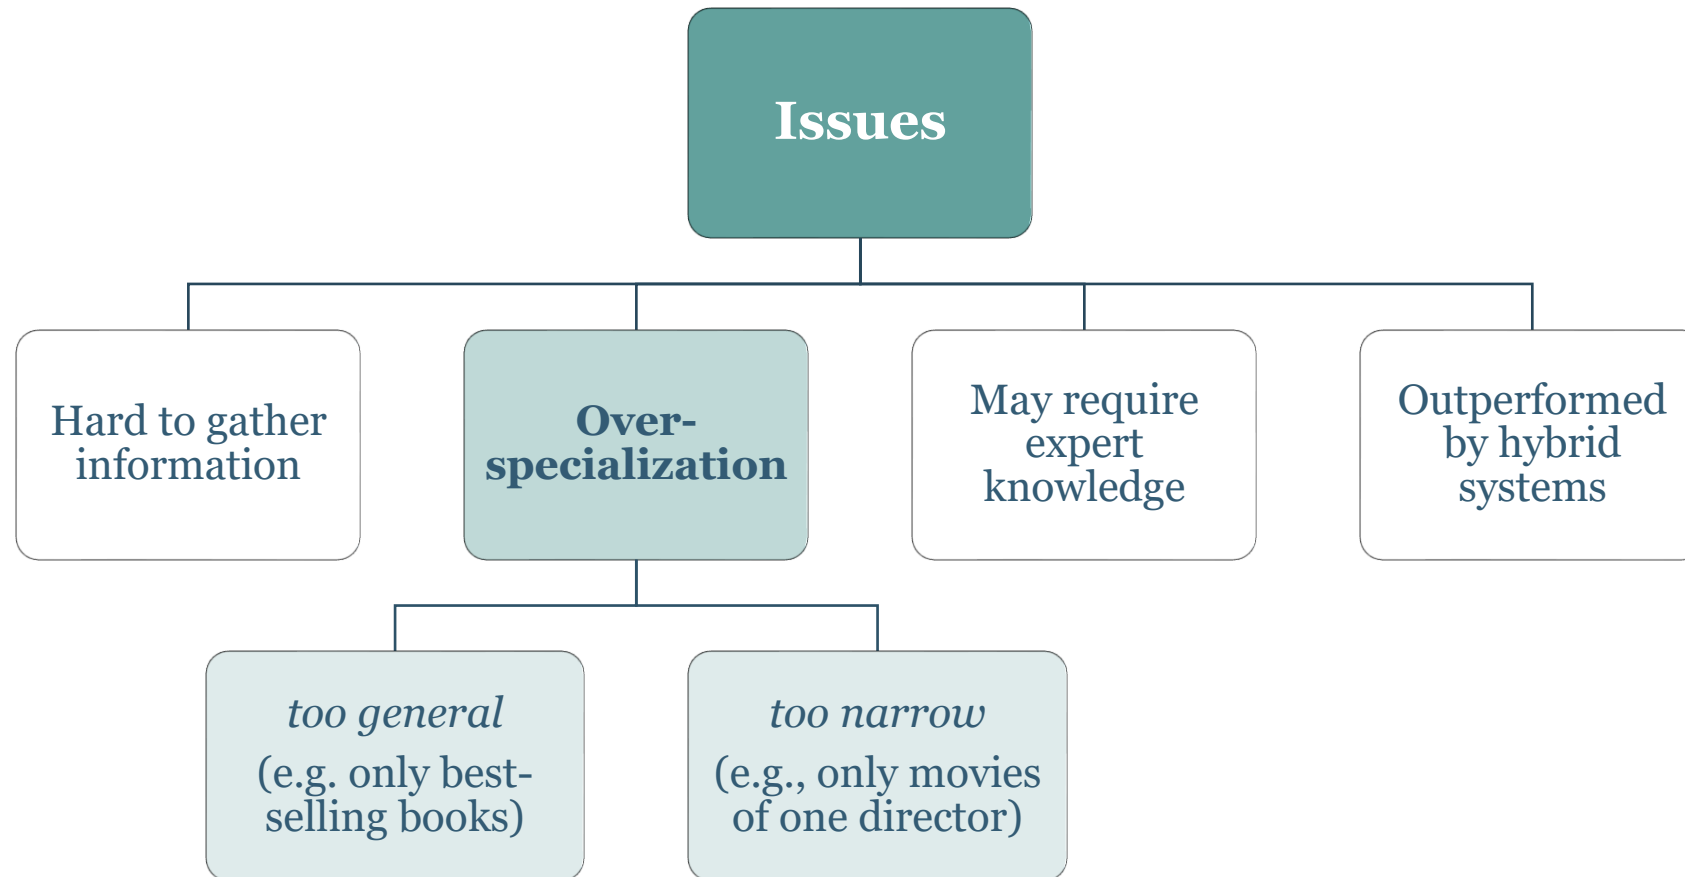
*Use standard definition of precision and recall.*

# Today's lecture

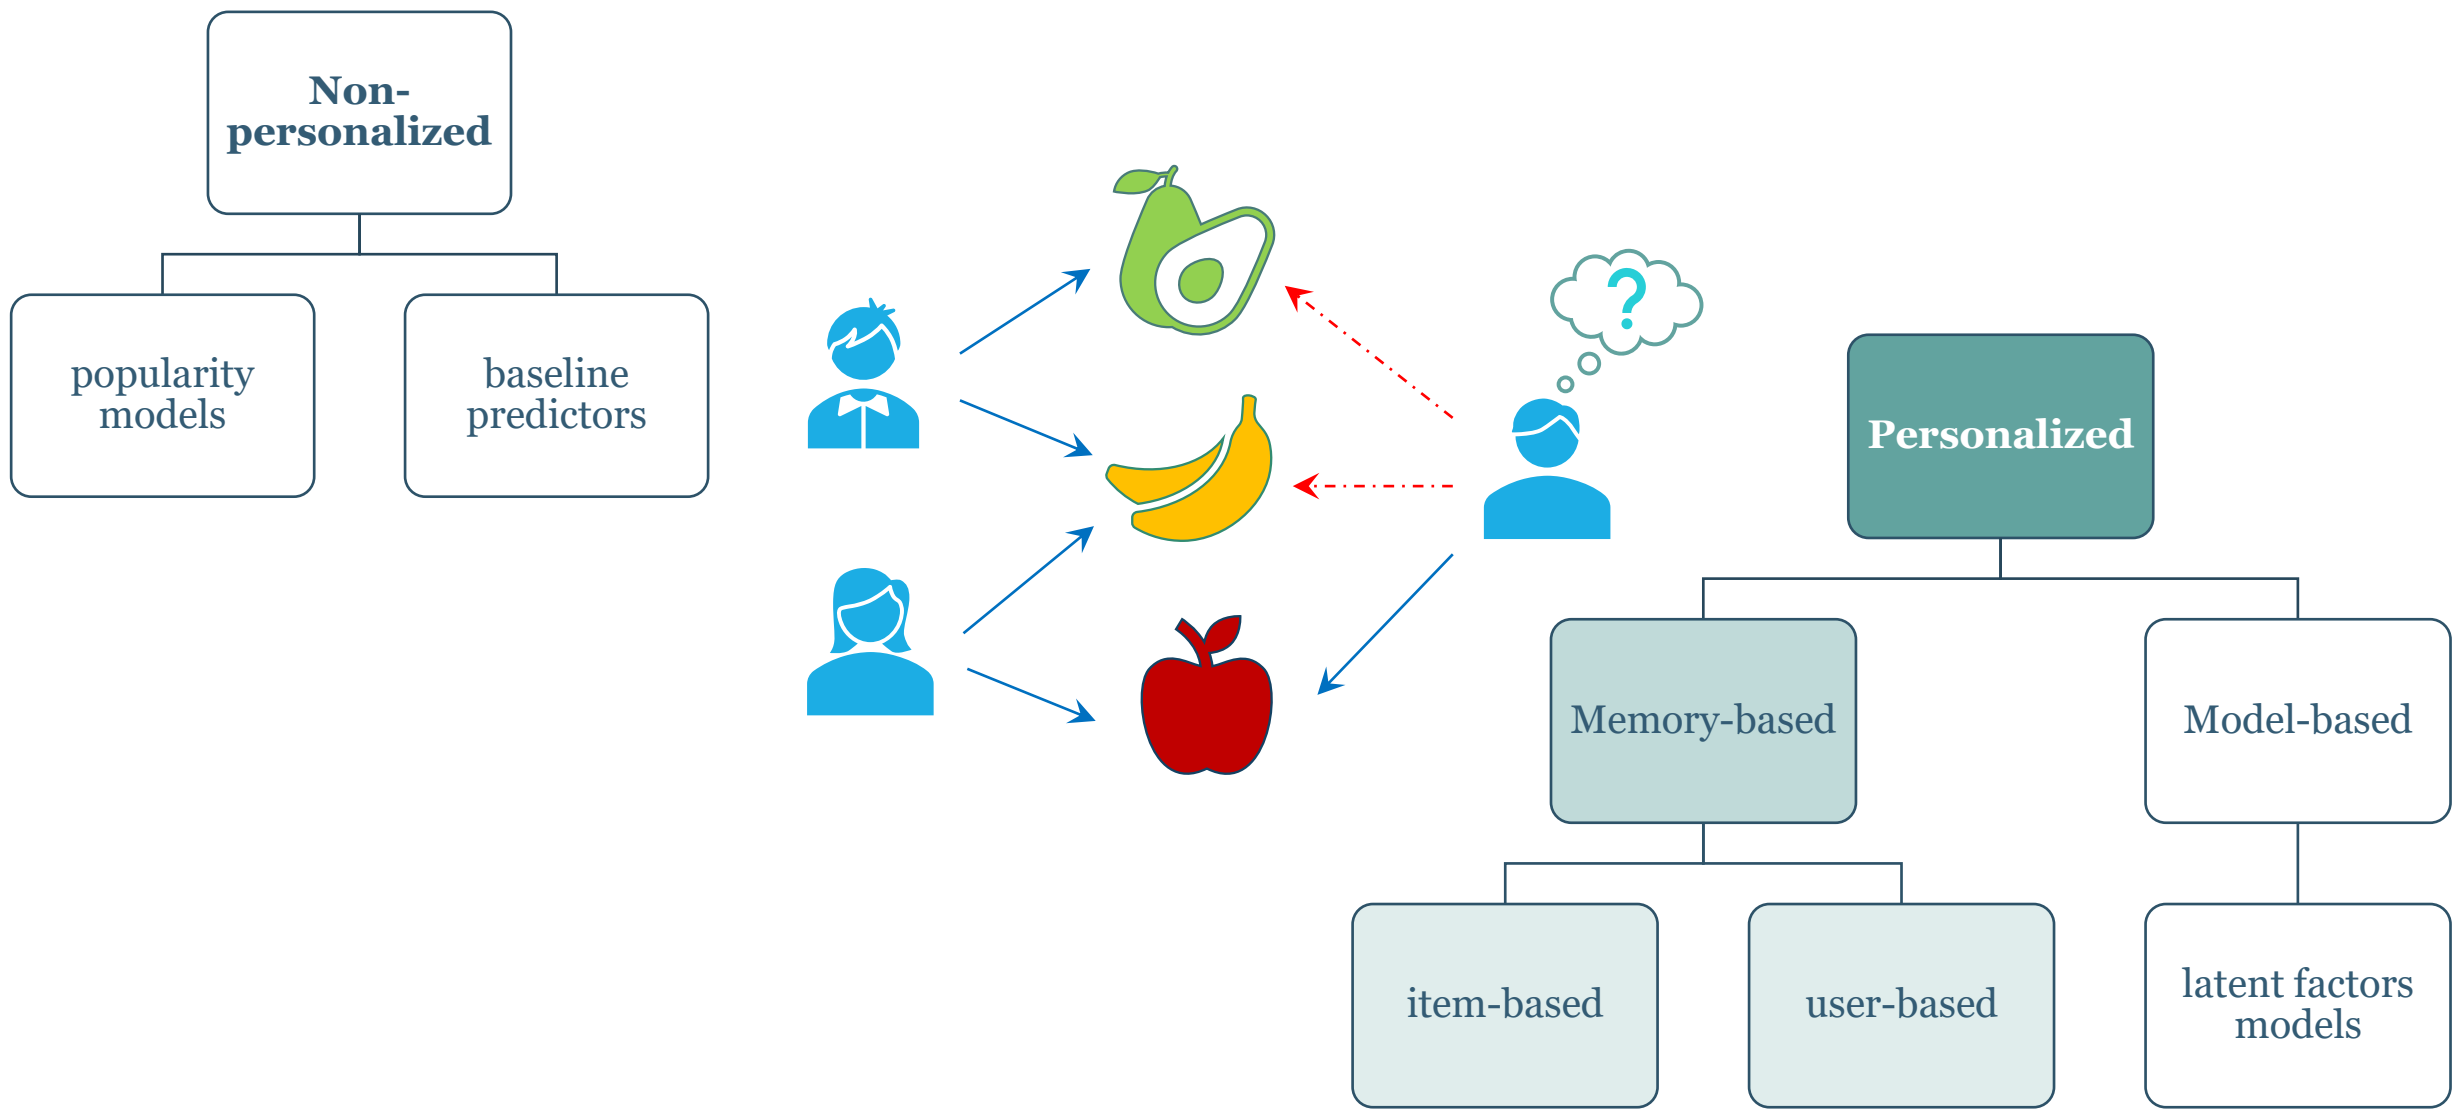
## Collaborative filtering

- memory-based approach
  - frequent pattern mining
  - nearest neighbors models

# Previously: content-based approach



# Collaborative Filtering: “wisdom of crowds”



# General workflow

Goal: predict user preferences based on prior user feedback and collective user behavior.

**collect data**

			
	?	?	3
	5	5	?
	4.5	?	4


user-movie matrix  $A$  of size  $M \times N$

$a_{ij}$  is a rating of  $i^{th}$  user for  $j^{th}$  movie

? - missing (unknown) values

**build model**

$f_U: \text{User} \times \text{Item} \rightarrow \text{Relevance}$



**generate recommendations**





# “Customers who like ... also like ...”



How do we implement that logic?


# Pure item-to-item (I2I)

## Typical transactions log:

user id	item id	transact.
0	575	view
0	1881	view
0	846	basket
1	1878	purchase
1	576	view
...	...	...



## Count co-occurrence of items:




$$\text{score}_{\text{I2I}}(u, i) = \sum_{\substack{j \in I_u \\ j \neq i}} \text{pairCount}(i, j)$$

Item-to-item is a strong baseline on very “sparse” datasets.

# Simplest item-to-item approach

## Convenient representation of logs– sparse matrix

user id	item id	transact.
0	575	view
0	1881	view
0	846	basket
1	1878	purchase
1	576	view



- Can be efficiently stored in CSR or CSC formats.
- Also enables efficient computations (especially useful for experiments).

# Computing I2I scores



- How to compute item-to-item co-occurrence matrix in symmetric case?
- How to compute similarity scores in that case?

$A$  — user-item interactions  $M \times N$

$$C = A^T A - \text{diag}(\text{diag}(A^T A))$$

$$p = [0 \underset{\uparrow}{1} \dots \underset{\uparrow}{1} \dots 0 \dots \underset{\uparrow}{1} \dots 0 \ 1]^T \in \mathbb{R}^N$$

$$z = Cp$$


# Computing I2I scores

$$C = A^T A - \text{diag}(\text{diag}(A^T A))$$

If  $p$  is a vector of known user preferences,  
then the vector of predicted relevance scores is:

$$r = Cp$$

Recommendations:

$$\text{toprec}(n) := \arg \max_j^n r_j$$


# Complexity analysis

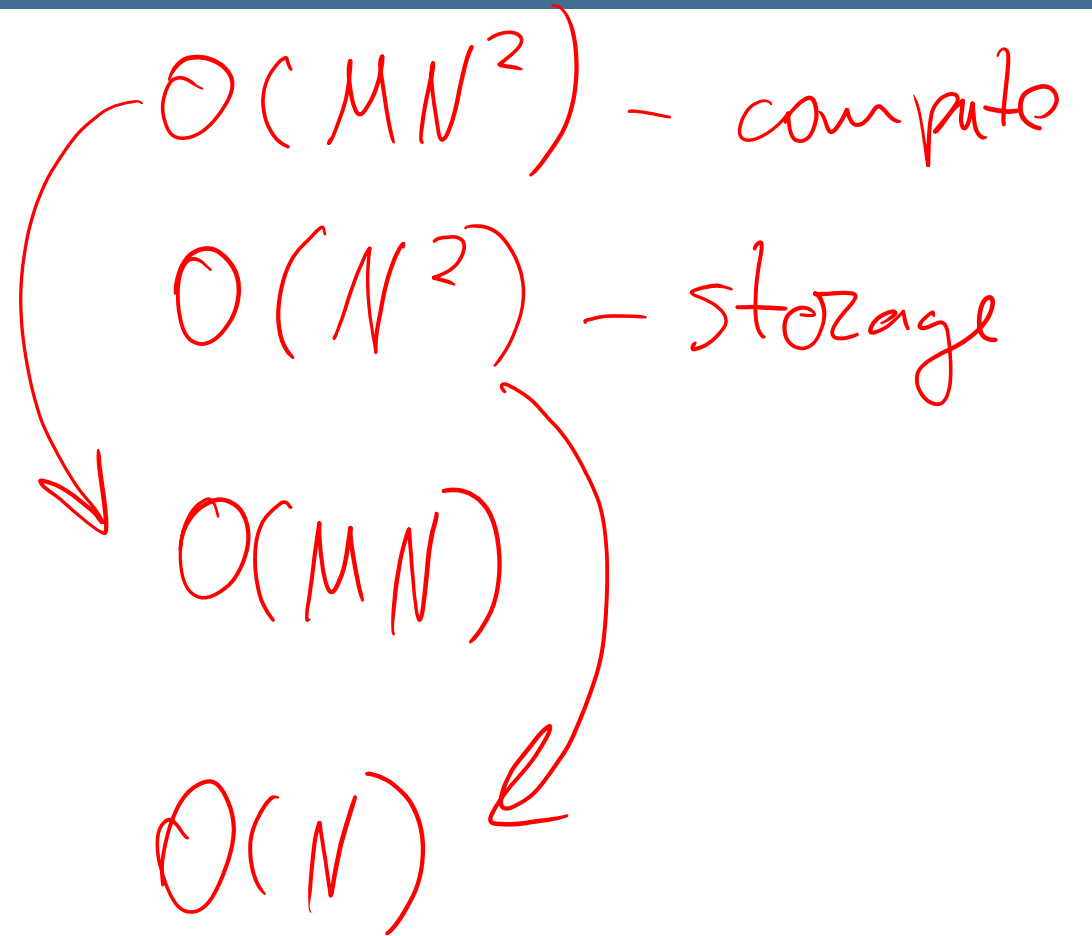
$$C = A^T A \quad M \times N$$

$$O(1)$$

$$O(\log N)$$

$$O(M N \log N)$$

$$O(N \log N)$$

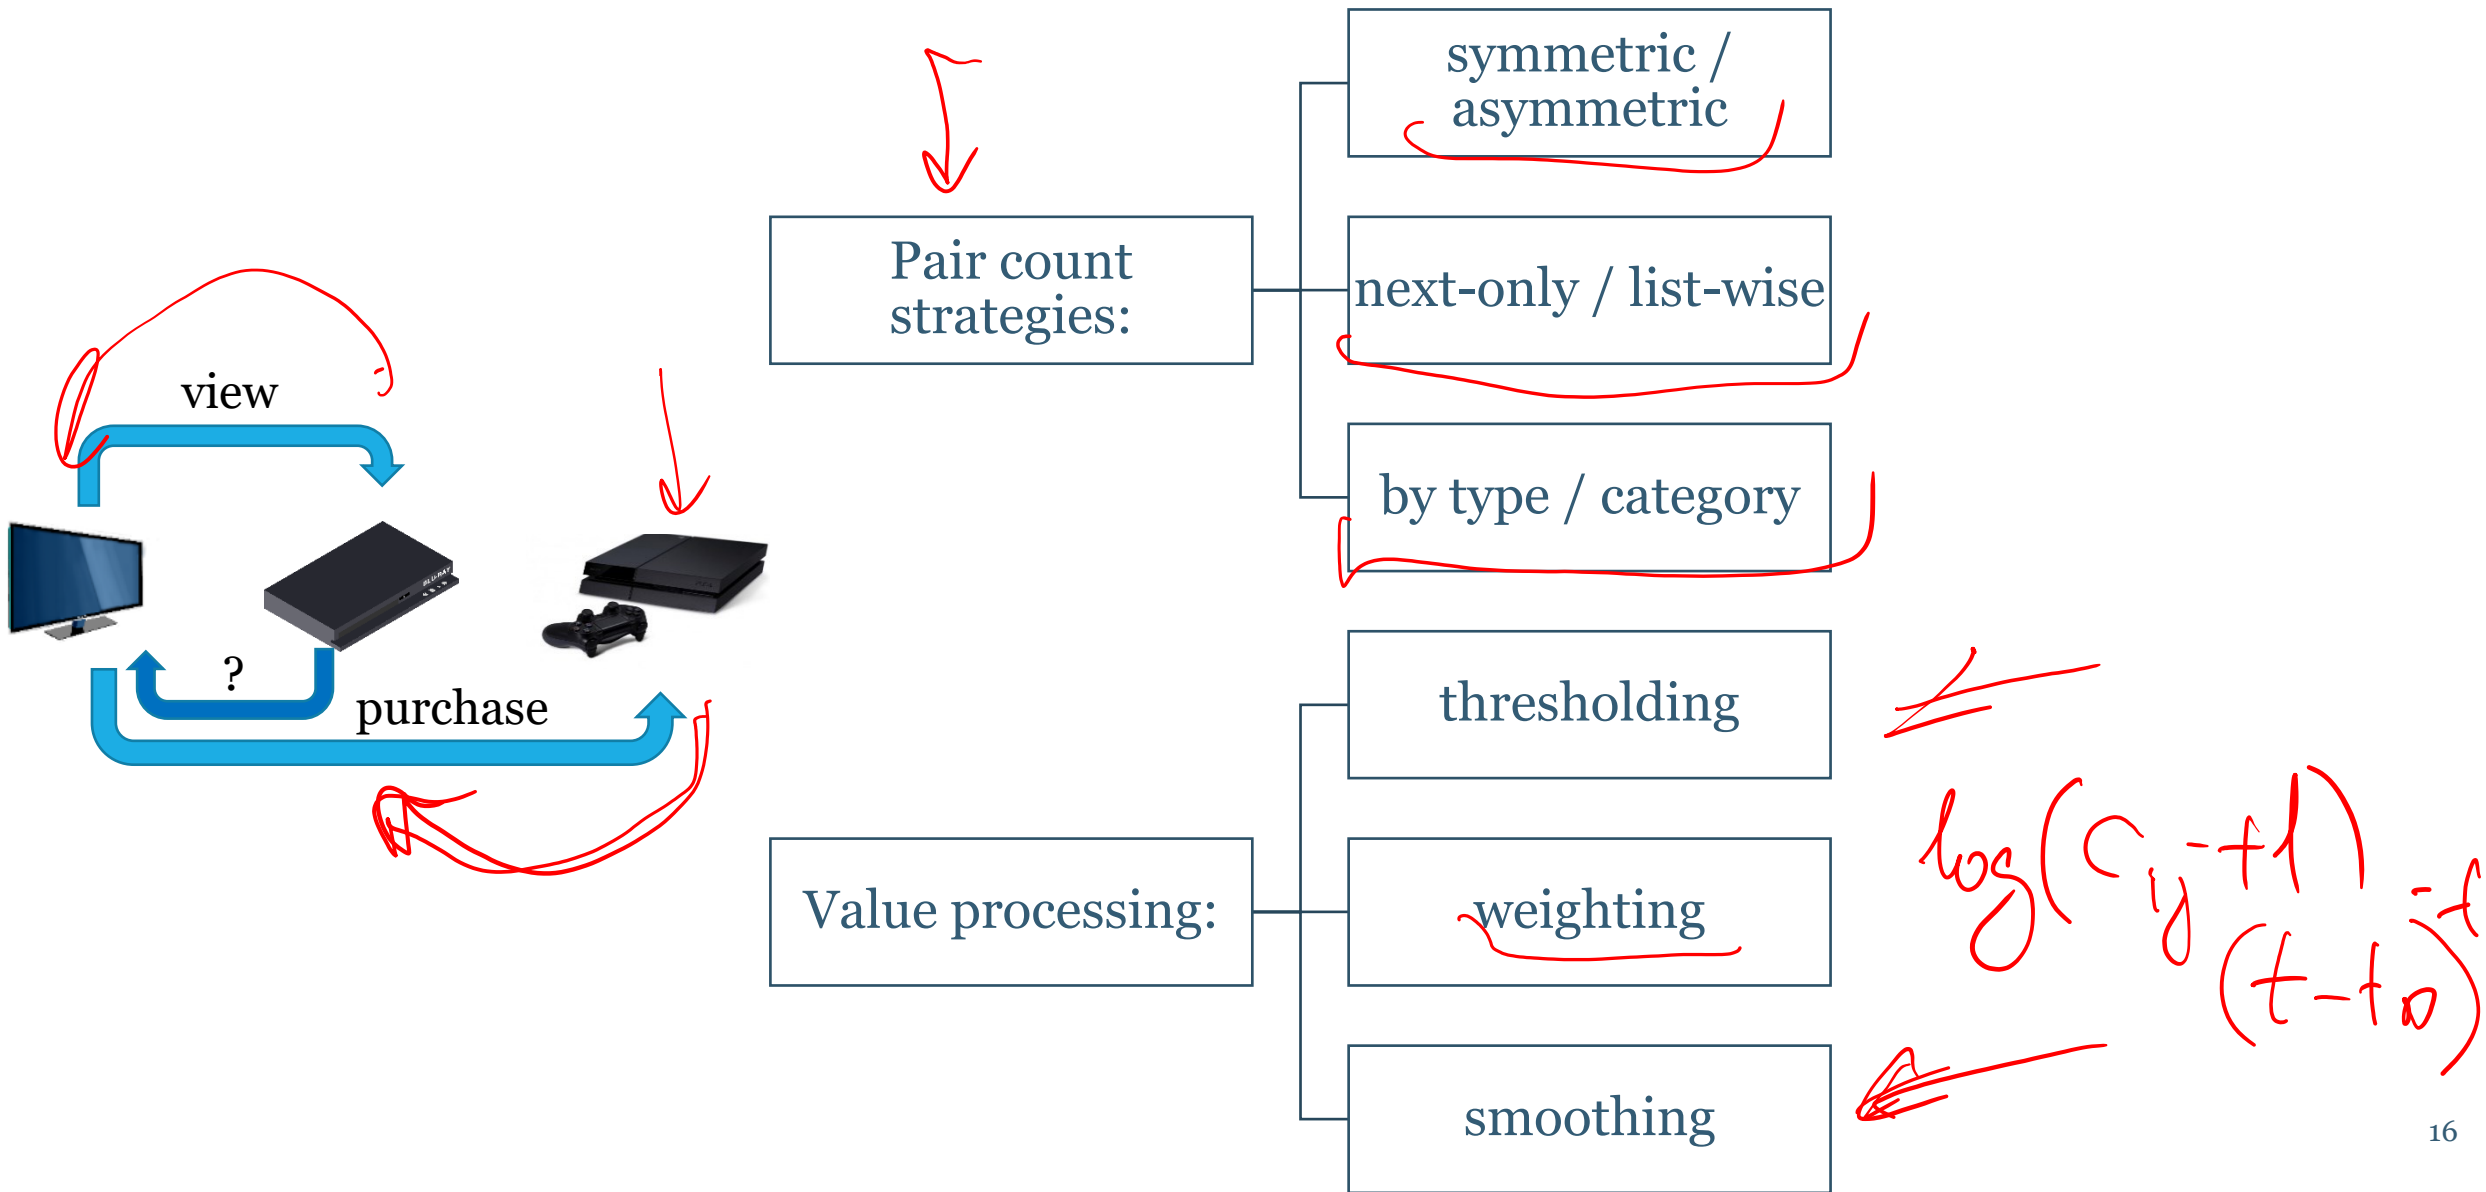


# Item-to-item issues

- somewhat obvious recommendations
  - high influence of popular items
- <sup>h</sup>i2i matrix can also become dense if there are too many interactions per user

$$C_{ij} \sim \vec{a}_i^T \vec{a}_j$$

# Item-to-item variants






# Case study: Amazon item-to-item

```

For each item in product catalog,  $I_1$ 
  For each customer  $C$  who purchased  $I_1$ 
    For each item  $I_2$  purchased by
      customer  $C$ 
      Record that a customer purchased  $I_1$ 
        and  $I_2$ 
    For each item  $I_2$ 
      Compute the similarity between  $I_1$  and  $I_2$ 
  
```


G. Linden, B. Smith and J. York, "Amazon.com recommendations: item-to-item collaborative filtering," in IEEE Internet Computing, vol. 7, no. 1, pp. 76-80, Jan.-Feb. 2003.

Total citations Cited by 8186



Iterative algorithm

Computes similarity of items based on user purchases.



$$\text{sim}(I_1, I_2) = \cos(\bar{a}_1, \bar{a}_2) = \frac{(\bar{a}_1, \bar{a}_2)}{\|\bar{a}_1\| \|\bar{a}_2\|}$$

$\bar{a}_k$  - "one-hot" representation of item  $k$

# Scalability trick: incremental updates in binary case

$$\text{sim}(i, j) = \frac{\bar{\mathbf{a}}_i^\top \bar{\mathbf{a}}_j}{\|\bar{\mathbf{a}}_i\| \cdot \|\bar{\mathbf{a}}_j\|} = \frac{\text{pairCount}(i, j)}{\sqrt{\text{itemCount}(i)} \cdot \sqrt{\text{itemCount}(j)}}, \quad j \neq i$$

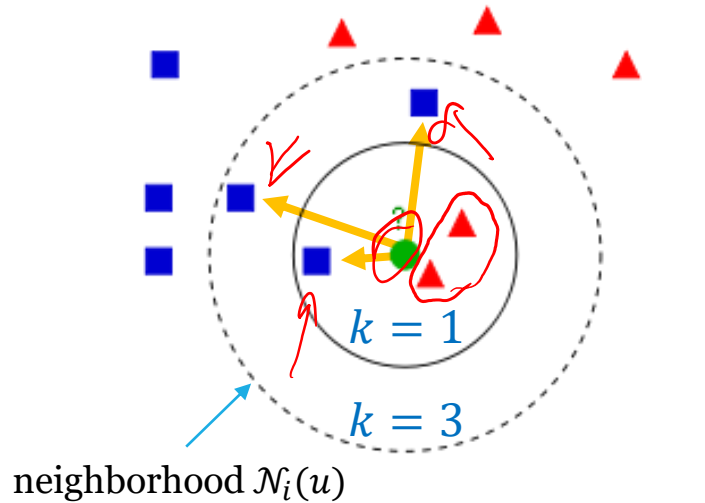
$$\|\bar{\mathbf{a}}_i\|^2 = \sum_u a_{ui}^2 = \sum_u a_{ui} = \text{itemCount}(i) \quad \bar{\mathbf{a}}_i^\top \bar{\mathbf{a}}_j = \sum_u a_{ui} a_{uj} = \text{pairCount}(i, j)$$

After observing  $\Delta A$  new interactions s.t.  $A' = A + \Delta A$ , the updated similarity is:

$$\text{sim}'(i, j) = \frac{\text{pairCount}(i, j) + \sum_u [\Delta A]_{ui} [\Delta A]_{uj}}{\sqrt{\text{itemCount}(i) + \sum_u [\Delta A]_{ui}} \cdot \sqrt{\text{itemCount}(j) + \sum_u [\Delta A]_{uj}}}$$

# Nearest neighbors models

# kNN-based approach



- user  $u$
- neighbors of user  $u$ , who rated item  $i$
- ▲ other users, who have not rated item  $i$

## User-based approach

- aggregated opinion of like-minded users:

$$\text{score}_{\text{uKNN}}(u, i) = \text{agg}_{v \in \mathcal{N}_i(u)} a_{vi}$$

## Item-based approach:

- aggregated rating of similarly rated items

$$\text{score}_{\text{iKNN}}(u, i) = \text{agg}_{j \in \mathcal{N}_u(i)} a_{uj}$$

# Simple user-based kNN

$$\text{score}_{\text{uKNN}}(u, i) = \frac{1}{|\mathcal{N}_i(u)|} \sum_{v \in \mathcal{N}_i(u)} a_{vj}$$



Potential issues:

- users may have very different interests
- neighborhood size is unlimited



# Improved user-based kNN

$$\text{score}_{\text{uKNN}}(u, i) = \frac{1}{z} \sum_{v \in \mathcal{N}_i(u)} \text{sim}(u, v) \cdot a_{vi}$$

$$\mathcal{N}_i(u) = U_i \setminus \{u\}, \quad z = \sum_{v \in \mathcal{N}_i(u)} |\text{sim}(u, v)|$$



Potential issues:

- ~~other users may have very different interests~~
- large neighborhood size

$$\sum_i w_i \alpha_i$$
$$w_i \in [0, 1]$$

# Dealing with large neighborhood size

- Storing similarities or on-the-fly computations?

- Aggressive subsampling

randomly  
sample from  $N_u(i)$   
recency-based  $N_{it}(u)$

- Approximate nearest neighbors

- e.g., NMSLib, Faiss, Annoy

- Dimensionality reduction

# Reducing neighborhood


- from  $N$  total entities sample  $n \ll N$
- select top- $k$  most similar among  $n$  samples,  $k \ll n$

Possible sampling strategies (must be fast):

- randomly
- most recent only
- most ratings in common (turns into fast MIPS problem)



# Example: local time-aware sampling



## Sampling strategy:

- select users that have items in common with a target user  $u$ 
  - each item of a neighbour-user must precede the corresponding item in the target user profile
  - filter out neighbours with too few items in common

## Additional weighting:

- users with no recent ratings  $\rightarrow$  lower weights
- active neighbour-user but old rating on a target item  $\rightarrow$  lower weights

	recent user ( $t_0 \approx t_{u'/1}$ )	old user ( $t_0 \gg t_{u'/1}$ )
recent item ( $t_{u'/1} \approx t_{u'/i}$ )	$\approx 0$	$t_0 - t_{u'/1}$
old item ( $t_{u'/1} \gg t_{u'/i}$ )	$t_{u'/1} - t_{u'/i}$	$t_0 - t_{u'/1}$

# Centered kNN

- baseline estimators contain most of useful signal
- every user may have individual “rating scale”

$$\text{score}_{\text{uKNN}}(u, i) = \bar{a}_u + \frac{1}{Z} \sum_{v \in \mathcal{N}_i(u)} \text{sim}(u, v) \cdot (a_{vi} - \bar{a}_v)$$

$\bar{a}_u$  - average rating of user  $u$

item-based kNN:

$$\text{score}_{\text{iKNN}}(u, i) = \bar{a}_i + \frac{1}{Z} \sum_{j \in \mathcal{N}_u(i)} \text{sim}(i, j) (a_{uj} - \bar{a}_j)$$

# Centered kNN

- baseline estimators contain most of useful signal
- every user may have individual “rating scale”

user-based kNN:

$$\text{score}_{\text{uKNN}}(u, i) = \bar{a}_u + \frac{1}{Z} \sum_{v \in \mathcal{N}_i(u)} \text{sim}(u, v) \cdot (a_{vi} - \bar{a}_v)$$

$\bar{a}_u$  - average rating of user  $u$

item-based kNN:

$$\text{score}_{\text{iKNN}}(u, i) = \bar{a}_i + \frac{1}{Z} \sum_{j \in \mathcal{N}_u(i)} \text{sim}(i, j) \cdot (a_{uj} - \bar{a}_j)$$

$\bar{a}_i$  - average rating of user  $i$

# Similarity measures

- Cosine Similarity
- Pearson Correlation
- Adjusted Cosine Similarity
- Jaccard Index
- Weighted Jaccard Index
- Asymmetric Similarities
- ...
- Spearman's Rank Correlation
- Kendall Tau



insensitive to ranking of “bad” items vs “good ” items

# Baseline-adjusted similarity

- **Pearson correlation** (adopted for CF):

$$\text{score}_{\text{Pearson}}(u, v) = \frac{\sum_{i \in I_u \cap I_v} (a_{ui} - \bar{a}_u)(a_{vi} - \bar{a}_v)}{\sqrt{\sum_{i \in I_u \cap I_v} (a_{ui} - \bar{a}_u)^2} \cdot \sqrt{\sum_{i \in I_u \cap I_v} (a_{vi} - \bar{a}_v)^2}}$$

$\bar{a}_u$  - average rating of user  $u$

- **Adjusted Cosine Similarity:**

$$\text{score}_{\text{AC}}(u, v) = \frac{\sum_{i \in I_u \cap I_v} (a_{ui} - \bar{a}_i)(a_{vi} - \bar{a}_i)}{\sqrt{\sum_{i \in I_u \cap I_v} (a_{ui} - \bar{a}_i)^2} \cdot \sqrt{\sum_{i \in I_u \cap I_v} (a_{vi} - \bar{a}_i)^2}}$$


$\bar{a}_i$  - average rating of item  $i$

# Jaccard Index

Item-based similarity:

$$\text{sim}_{\text{JI}}(i, j) = \frac{|U_i \cap U_j|}{|U_i \cup U_j|}$$

$$|U_i \cup U_j| = |U_i| + |U_j| - |U_i \cap U_j|$$




# Weighted Jaccard Index


- Jaccard Index only operates on sets
- often some values are associated with interactions (e.g., ratings, frequencies)

$$\text{sim}_{\text{WJI}}(u, v) = \frac{\sum_{i=1}^N \min\{w_i(a_u), w_i(a_v)\}}{\sum_{i=1}^N \max\{w_i(a_u), w_i(a_v)\}},$$

$$w_i(a_u) = f(a_{ui})$$



Images from: <https://moultano.wordpress.com/2018/11/08/minhashing-3kbzhxsxyg4467-6/>



• intersection = [ banana, avocado ]

• union = [ banana, banana, banana, avocado, avocado, apple, apple ]

•  $\text{score}_{\text{WJI}}(\text{user1}, \text{user2}) = \frac{2}{7}$

# kNN in matrix form

Element-wise weighting for user-based KNN:

$$r_{ui} = \frac{1}{z_{ui}} \cdot \sum_{v \in \mathcal{N}_i(u)} \text{sim}(u, v) \cdot a_{vi}$$

Handwritten annotations for the equation above:  
 - An arrow points to  $r_{ui}$ .  
 - An arrow points to the denominator  $z_{ui}$ .  
 - An arrow points to the summation index  $v \in \mathcal{N}_i(u)$ .  
 - A bracket groups the entire summation term  $\sum_{v \in \mathcal{N}_i(u)} \text{sim}(u, v) \cdot a_{vi}$ .  
 - Inside the bracket, an arrow points to  $\text{sim}(u, v)$  and another points to  $a_{vi}$ .

$$\mathcal{N}_i(u) = U_i \setminus \{u\}$$

we impute  $A$  with 0's

$$R = K A \odot K B$$

Handwritten annotation: "elementwise division" with an arrow pointing to the  $\odot$  operator.

$$z_{ui} = \sum_{v \in \mathcal{N}_i(u)} |\text{sim}(u, v)|$$

Handwritten annotation: An arrow points to the summation index  $v \in \mathcal{N}_i(u)$ .

$$[B]_{ij} = \begin{cases} 1 & a_{ij} \text{ known} \\ 0 & \text{otherwise} \end{cases}$$

Handwritten annotation: An arrow points to the  $[B]$  term.

$M \times M$  zero-diagonal.  
 $K$  - similarity matrix  
 $A \in M \times N$



# kNN in matrix form

Row-wise weighting for user-based KNN:

$$r_{ui} = \frac{1}{z_{ui}} \cdot \sum_{v \in \mathcal{N}_i(u)} \underbrace{\text{sim}(u, v) \cdot a_{vi}}_{\substack{\mathcal{N}_i(u) = U \setminus \{u\} \\ \text{i.e., "explicit" 0's}}}$$

$$z_{ui} = \sum_{v \in U \setminus \{u\}} |\text{sim}(u, v)| \quad [D]_{ui} = z_{ui}$$

$$R = D^{-1} K A$$

# kNN weighting schemes

$K$  – user similarity,  $k_{ii} = 0, k_{ij} \geq 0, i \neq j$ ;  $S$  – item similarity matrix,  $s_{ii} = 0, s_{ij} \geq 0, i \neq j$ .

## element-wise weighting:

- User-based:


$$R = KA \oslash (KB)$$

$$b_{ui} = \begin{cases} 1, & \text{if } a_{ui} \text{ is known} \\ 0 & \text{otherwise} \end{cases}$$

- Item-based:

$$R = AS^T \oslash (BS^T)$$

- filters known ratings only
- better for rating prediction



## row-wise / no weighting:

- User-based:

$$R = D_K^{-1}KA$$

$$D_K = \text{diag}(K\mathbf{e}) \text{ or } D_K = I$$

$$\mathbf{e} = [1, 1, \dots, 1]^T$$

- Item-based

$$R = AS^T D_S^{-1}$$

$$D_S = \text{diag}(S\mathbf{e}) \text{ or } D_S = I$$

- assumes 0-imputation of unknowns
- better for top- $n$  recommendations

Let's implement simple KNN  
models

# User-based kNN for top-n recommendations

- Row-wise weighting:

$$R = D^{-1}KA$$

$$r_{ui} = w_u \cdot \sum_{v \in \mathcal{N}(u)} \text{sim}(u, v) \cdot a_{vi}$$


- Is it different from the unweighted case?

- Alternative (column-wise) weighting:

$$R = KD^{-1}A$$

$$r_{ui} = \sum_{v \in \mathcal{N}(u)} \text{sim}(u, v) \cdot w_v \cdot a_{vi}$$

# kNN with asymmetric similarity

kNN similarity (e.g., item-based): 

row-wise weighted symmetric  $\rightarrow$  unweighted asymmetric

$$S_{\text{asym}} = D^{-\alpha} S$$

$$R = A S_{\text{asym}}^T$$

$$S = S^T$$
$$S_{\text{asym}} \neq S_{\text{asym}}^T$$

Example: cosine similarity, assuming  $d_{ii} = \|\bar{\mathbf{a}}_i\|$ :

$$\text{sim}(i, j) = [S_{\text{asym}}]_{ij} = \frac{\bar{\mathbf{a}}_i^T \bar{\mathbf{a}}_j}{\|\bar{\mathbf{a}}_i\|^{1+\alpha} \cdot \|\bar{\mathbf{a}}_j\|}$$

For binary data,  $\alpha = -1$  gives a simple conditional probability  $p(i|j)$

# Popularity effect in asymmetric similarity

$$\text{sim}(i, j) = [D^{-\alpha} S]_{ij} = \frac{\bar{\mathbf{a}}_i^\top \bar{\mathbf{a}}_j}{\|\bar{\mathbf{a}}_i\|^{1+\alpha} \cdot \|\bar{\mathbf{a}}_j\|}$$

## What do we recommend?

$$\begin{array}{ccccc} \text{popular item} & \xrightarrow{\text{sim}} & \text{unpopular} & \text{vs.} & \text{unpopular} \xrightarrow{\text{sim}} \text{popular item} \\ i & & j & & i \qquad j \end{array}$$

# Popularity effect in asymmetric similarity

popular item  $\xrightarrow{\text{sim}}$  unpopular      vs.      unpopular  $\xrightarrow{\text{sim}}$  popular item

$\alpha < 0$        $\alpha > 0$

- Popular products  $\rightarrow$  too trivial recommendations.
  - Easy to guess but low value for users + low diversity.
  - Recommending niche products increases diversity.
  - For users with generic tastes may not fit well.
- 
- Observation: popular items are not very descriptive of users interests.
  - Suggest a normalization that would improve item-KNN recommendations.

New scheme – emphasizing contribution of specific user tastes:

$$S_{\text{asym}} = SD^{-\beta}, \quad \text{sim}(i, j) = \frac{\bar{\mathbf{a}}_i^\top \bar{\mathbf{a}}_j}{\|\bar{\mathbf{a}}_i\| \cdot \|\bar{\mathbf{a}}_j\|^{1+\beta}}$$

Let's implement new weighting



# Useful reading on similarity in kNN

Herlocker, Jon, Joseph A. Konstan, and John Riedl. "An empirical analysis of design choices in neighborhood-based collaborative filtering algorithms." *Information retrieval* 5 (2002): 287-310.

Deshpande, Mukund, and George Karypis. "Item-based top-n recommendation algorithms." *ACM Transactions on Information Systems (TOIS)* 22, no. 1 (2004): 143-177.

Koenigstein, Noam, and Yehuda Koren. "Towards scalable and accurate item-oriented recommendations." In *Proceedings of the 7th ACM conference on Recommender systems*, pp. 419-422. 2013.

Naumov, Sergey, Marina Ananyeva, Oleg Lashinin, Sergey Kolesnikov, and Dmitry I. Ignatov. "Time-Dependent Next-Basket Recommendations." In *European Conference on Information Retrieval*, pp. 502-511. Cham: Springer Nature Switzerland, 2023.

# Summary

## Key advantages:


- easy to implement
- intuitive explanations
- good baseline
- suitable for very sparse datasets

## Scalability:

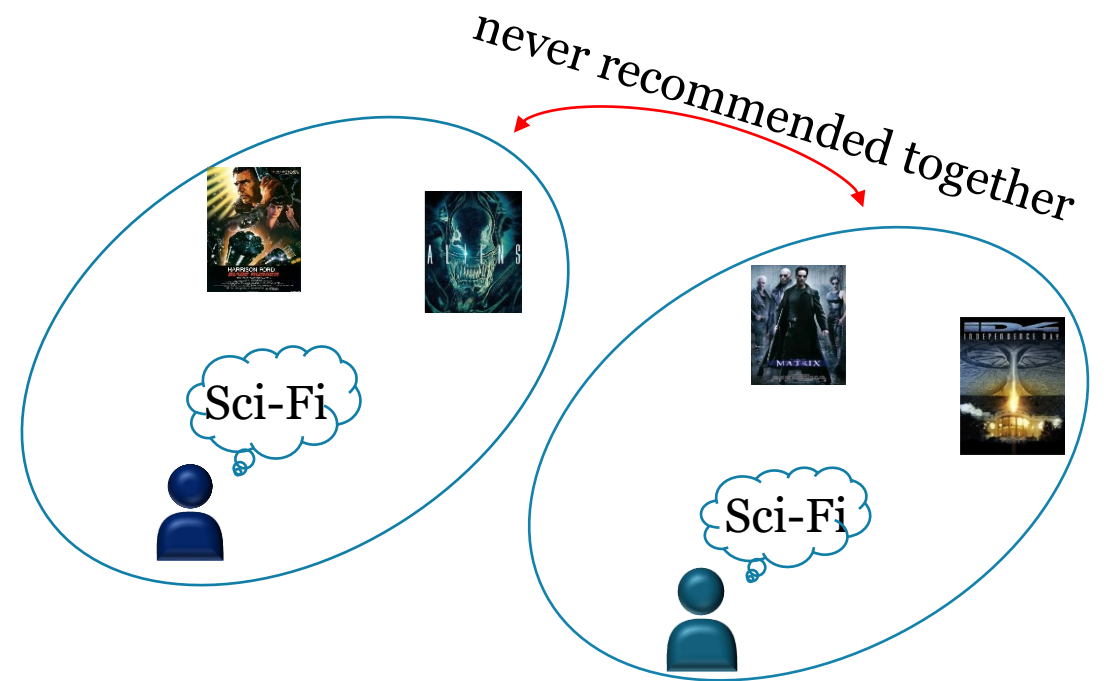
- computational complexity in the worst case:  $O(MN^2)$  or  $O(NM^2)$
- storage complexity in the worst case:  $O(M^2)$  or  $O(N^2)$
- due to sparsity real complexity can be significantly decreased
- could additionally limit the number of neighbors
- sometimes incremental updates are possible

# Limited coverage problems

## Unreliable correlations



## Weak generalization



# When to use user-based vs item-based?

- Depends on # users and # items
- Depends on dynamics
- Item-based recommendations are easier to explain.
- User-based recommendations increase serendipity.