

Recommender Systems

Lecture 4

Previous lecture

Lifecycle of a recsys experiment

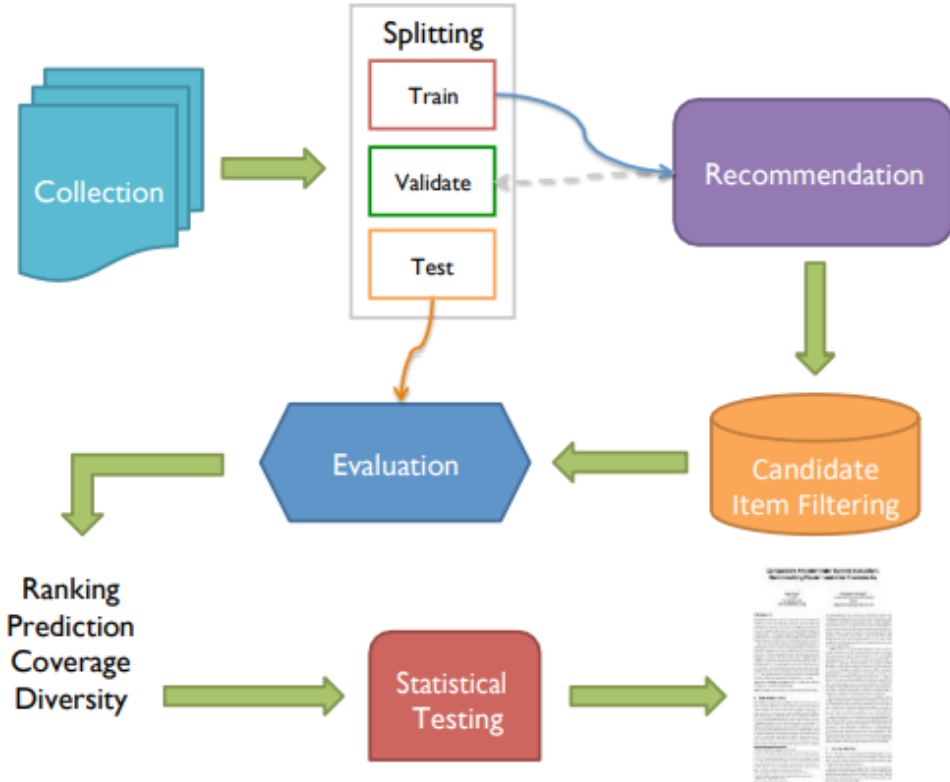



Image source: Bellogín, Alejandro, and Alan Said. "Improving accountability in recommender systems research through reproducibility." User Modeling and User-Adapted Interaction (2021): 1-37.

Data splitting



Splitting options:

- entry- or user-wise
- warm start
 - strong/weak generalization

Holdout sampling

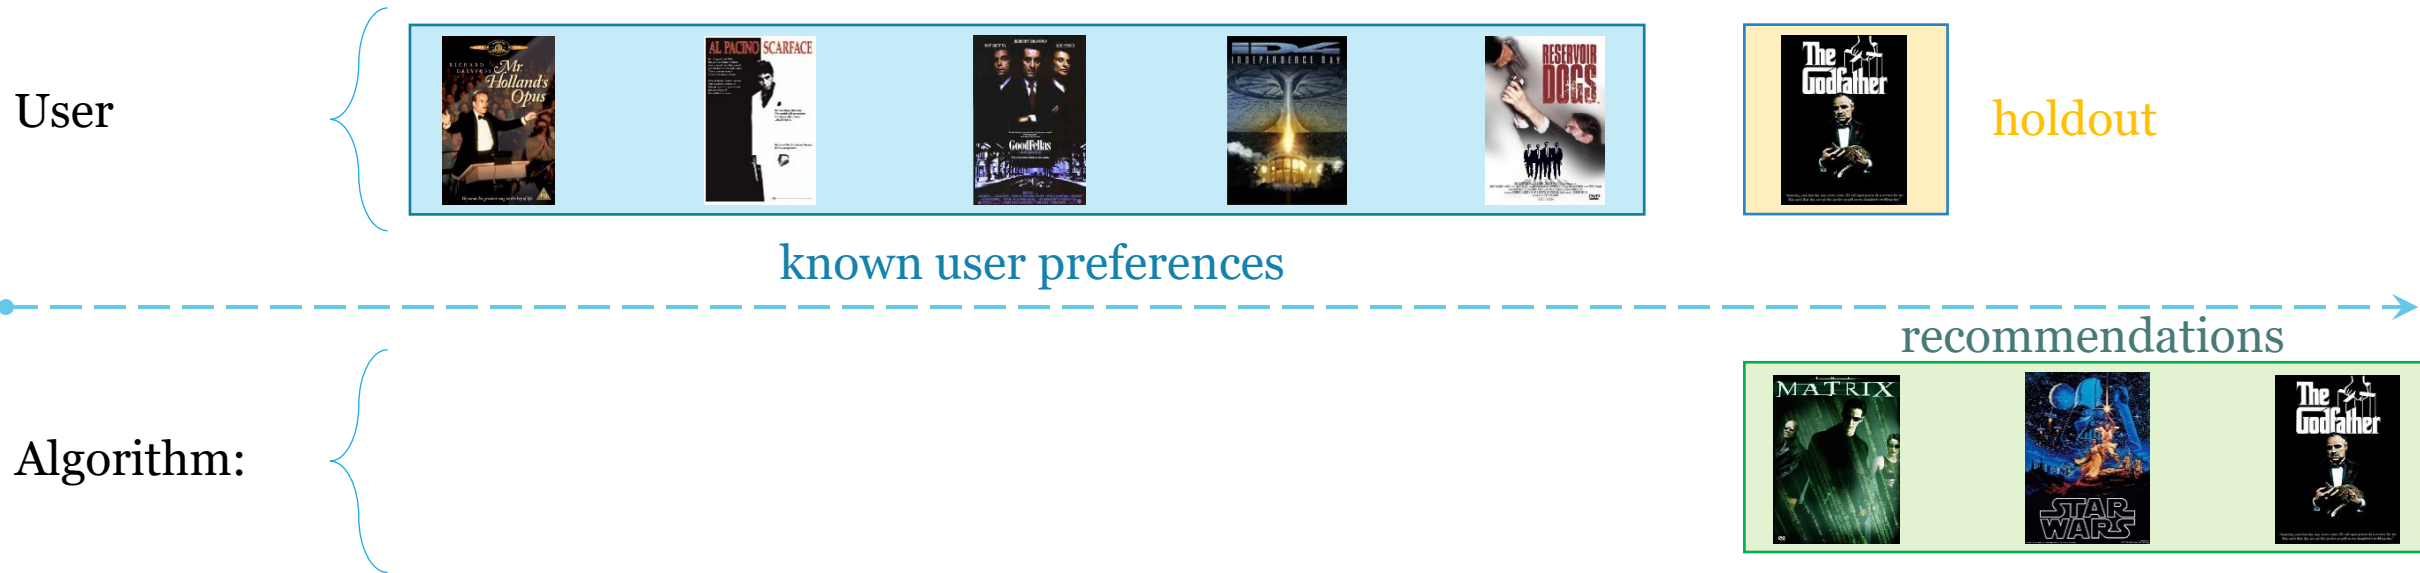
Strategies

- Random
- Rating-based (e.g., top-rated)
- Temporal (e.g., most recent)

Sample size

- Fixed number of items (e.g., 1)
- Fixed percentage of items

Previous lecture



$$HR = \frac{1}{\#(\text{test users})} \sum_{\text{test users}} \text{hit}$$

$$\text{hit} = \begin{cases} 1 & \text{if holdout item is in recommended items,} \\ 0 & \text{otherwise.} \end{cases}$$

$$MRR = \frac{1}{\#(\text{test users})} \sum_{\text{test users}} \frac{1}{\text{hit rank}}$$

hit rank = position of the item in the recommendations list

Typically computed: metric@n, where n = #recommended items, e.g. Recall@n, MRR@n, etc.

Task

Consider top-2 recommender for 10 users from 20-items (2 items per user).

What's better:

1) Correctly recommend 2 items to each of 5 users?

or

2) 1 item to each of 10 users?

Why?

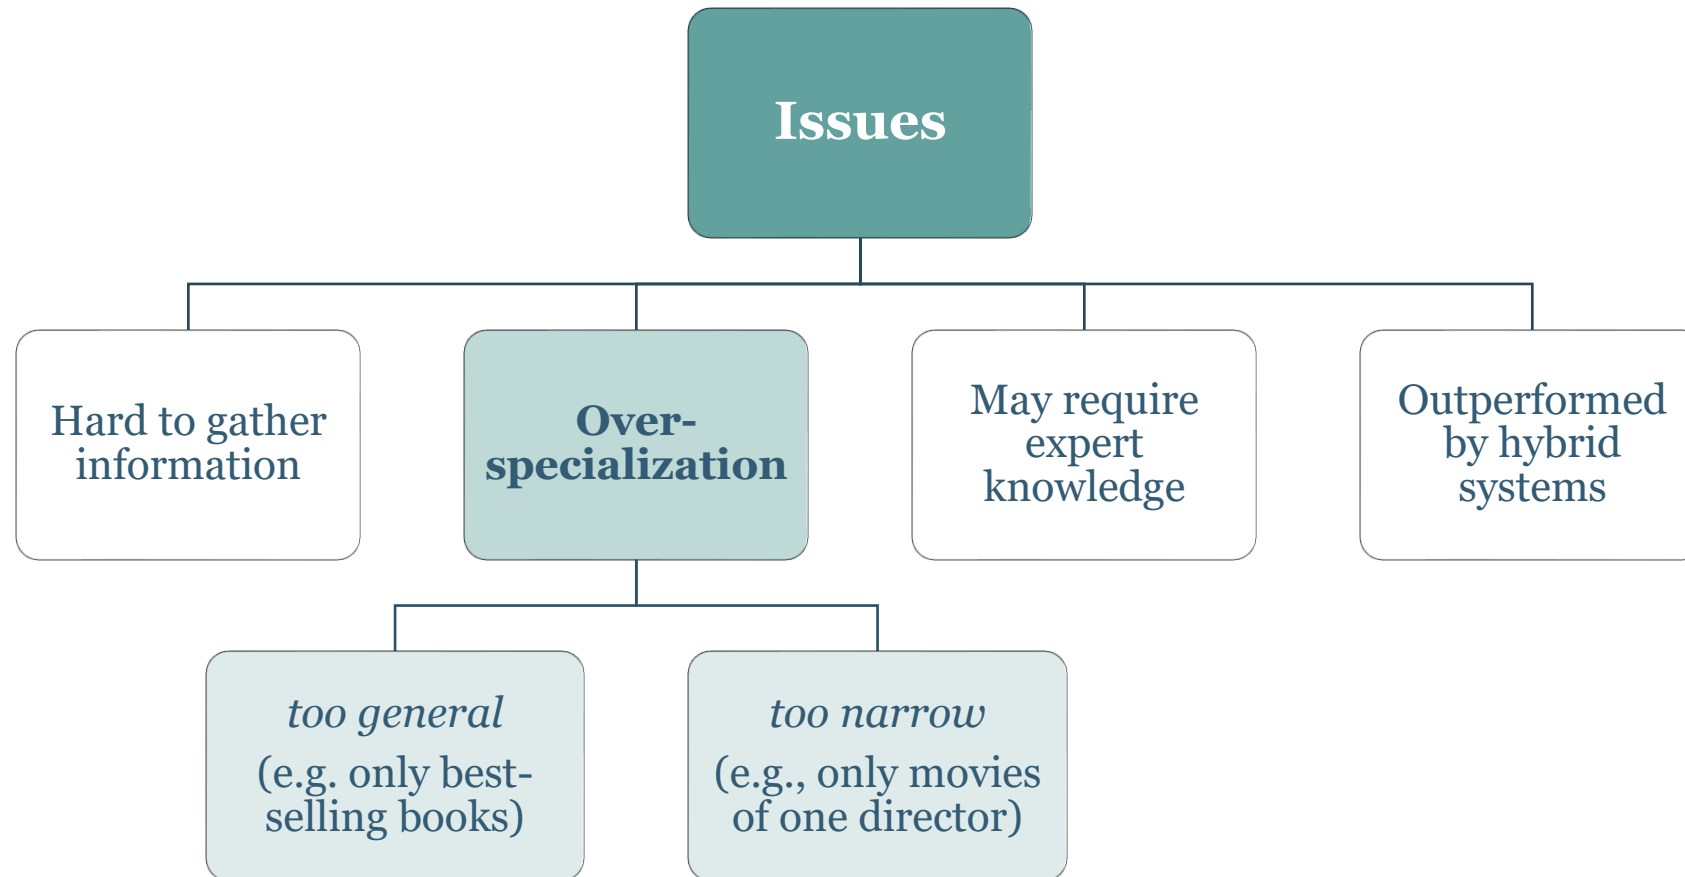
Use standard definition of precision and recall.

Today's lecture

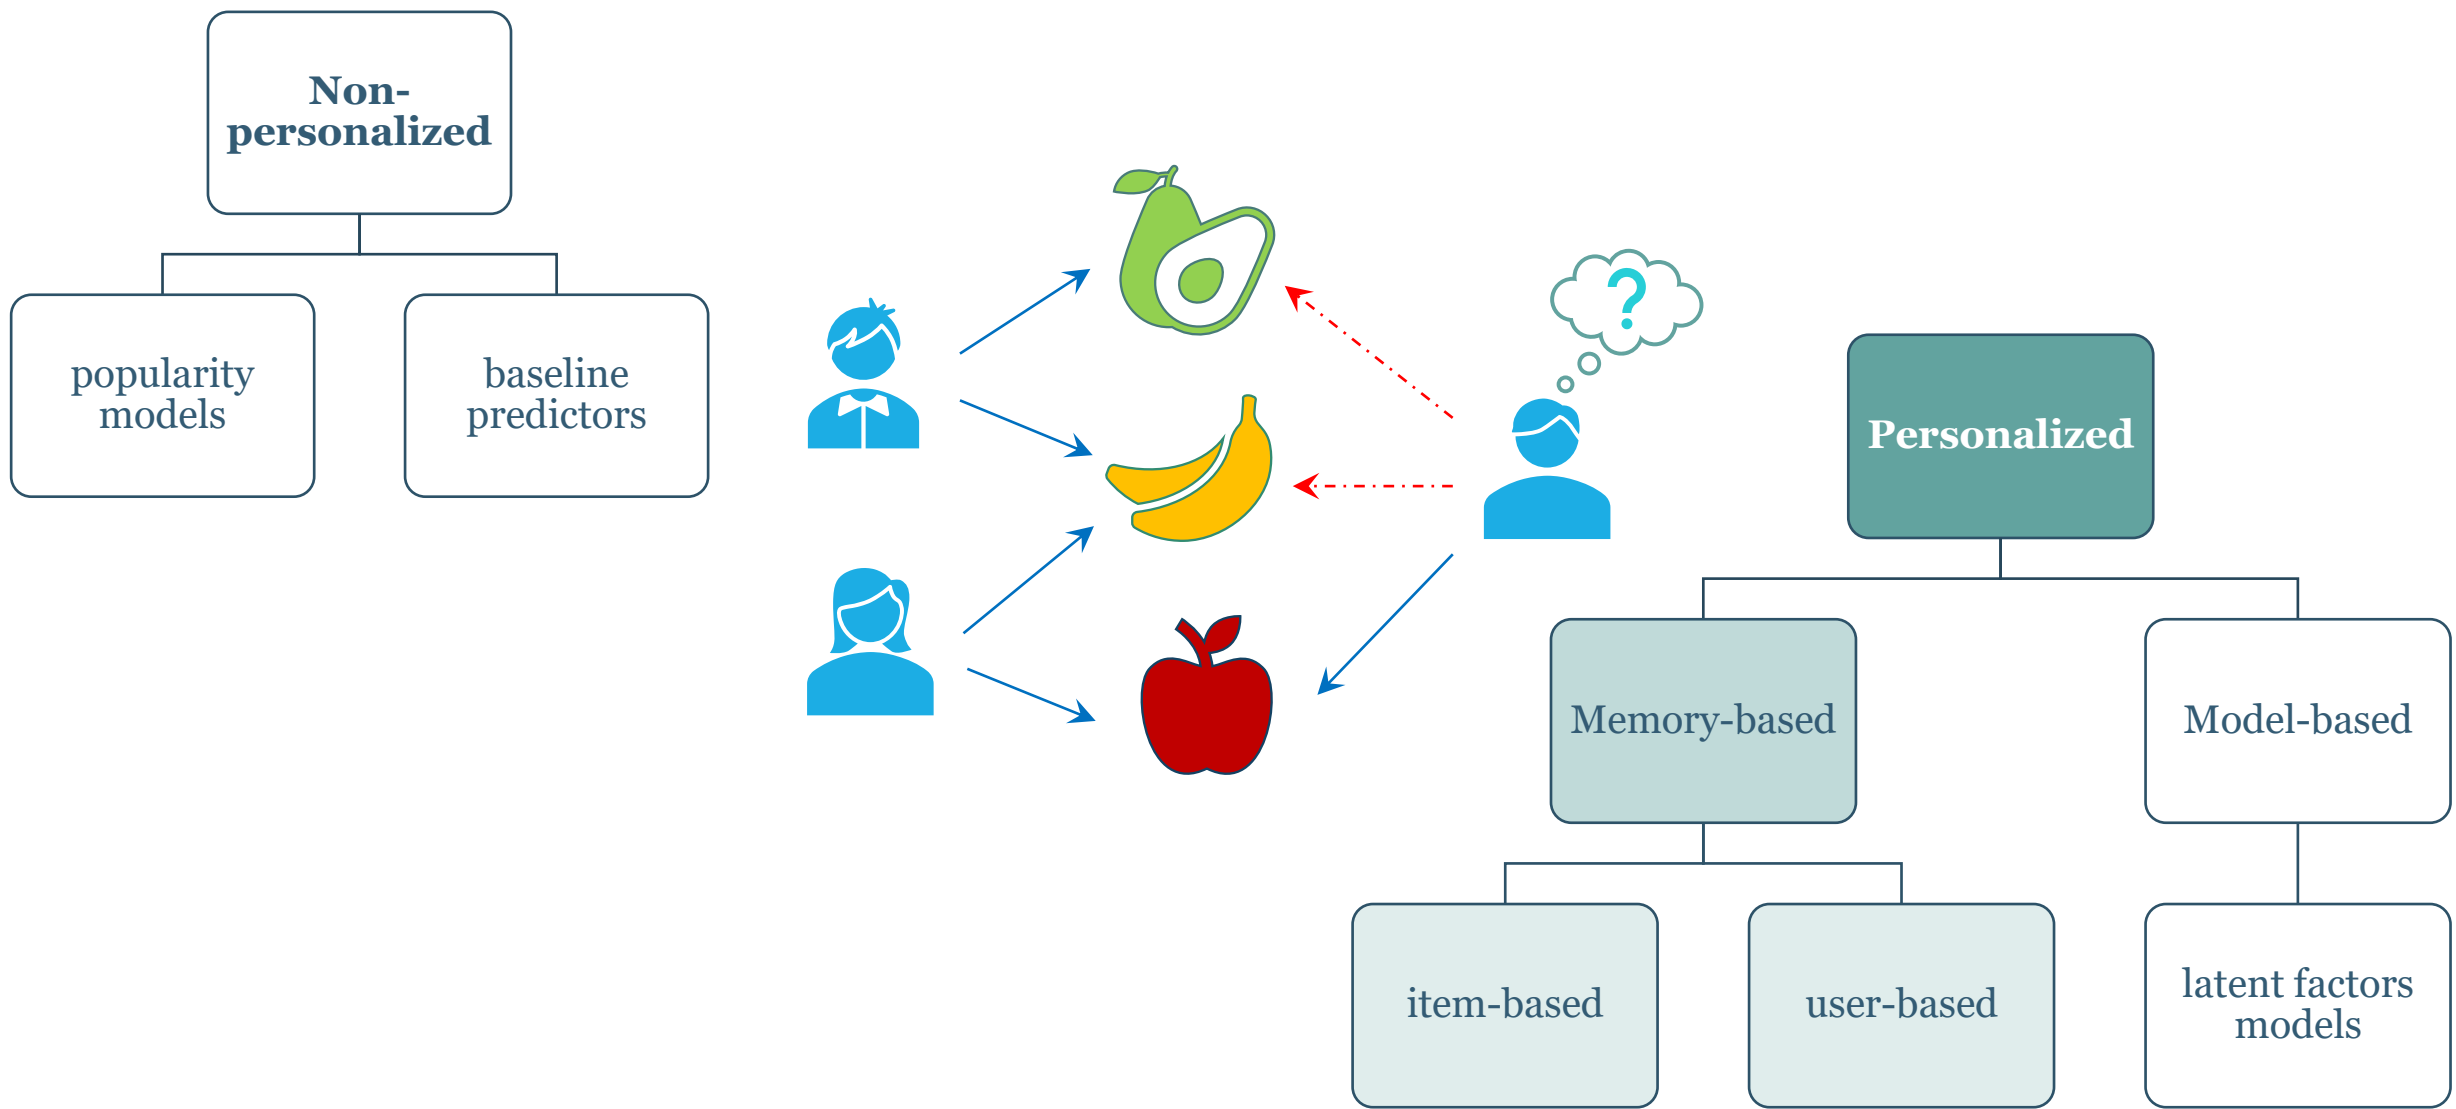
Collaborative filtering

- memory-based approach
 - frequent pattern mining
 - nearest neighbors models

Previously: content-based approach



Collaborative Filtering: “wisdom of crowds”



General workflow

Goal: predict user preferences based on prior user feedback and collective user behavior.

collect data

			
	?	?	3
	5	5	?
	4.5	?	4


user-movie matrix A of size $M \times N$

a_{ij} is a rating of i^{th} user for j^{th} movie


? - missing (unknown) values

build model

$f_U: \text{User} \times \text{Item} \rightarrow \text{Relevance}$



generate recommendations



“Customers who like ... also like ...”



How do we implement that logic?


Pure item-to-item (I2I)

Typical transactions log:

user id	item id	transact.
0	575	view
0	1881	view
0	846	basket
1	1878	purchase
1	576	view
...



Count co-occurrence of items:




$$\text{score}_{\text{I2I}}(u, i) = \sum_{\substack{j \in I_u \\ j \neq i}} \text{pairCount}(i, j)$$

Item-to-item is a strong baseline on very “sparse” datasets.

Simplest item-to-item approach

Convenient representation of logs– sparse matrix

user id	item id	transact.
0	575	view
0	1881	view
0	846	basket
1	1878	purchase
1	576	view



- Can be efficiently stored in CSR or CSC formats.
- Also enables efficient computations (especially useful for experiments).

Computing I2I scores



- How to compute item-to-item co-occurrence matrix in symmetric case?
- How to compute similarity scores in that case?

Computing I2I scores

$$C = A^T A - \text{diag}(\text{diag}(A^T A))$$

If p is a vector of known user preferences,
then the vector of predicted relevance scores is:

$$r = Cp$$

Recommendations:


$$\text{toprec}(n) := \arg \max_j^n r_j$$

Complexity analysis

Item-to-item issues

- somewhat obvious recommendations
 - high influence of popular items
- $i2i$ matrix can also become dense if there are too many interactions per user

Item-to-item variants



Pair count
strategies:

symmetric /
asymmetric

next-only / list-wise

by type / category

Value processing:

thresholding

weighting

smoothing


Case study: Amazon item-to-item

```

For each item in product catalog,  $I_1$ 
  For each customer  $C$  who purchased  $I_1$ 
    For each item  $I_2$  purchased by
      customer  $C$ 
      Record that a customer purchased  $I_1$ 
        and  $I_2$ 
  For each item  $I_2$ 
    Compute the similarity between  $I_1$  and  $I_2$ 
    
```


G. Linden, B. Smith and J. York, "Amazon.com recommendations: item-to-item collaborative filtering," in IEEE Internet Computing, vol. 7, no. 1, pp. 76-80, Jan.-Feb. 2003.

Total citations Cited by 8186



Iterative algorithm

Computes similarity of items based on user purchases.



$$\text{sim}(I_1, I_2) = \cos(\bar{a}_1, \bar{a}_2) = \frac{(\bar{a}_1, \bar{a}_2)}{\|\bar{a}_1\| \|\bar{a}_2\|}$$

\bar{a}_k - "one-hot" representation of item k

Scalability trick: incremental updates in binary case

$$\text{sim}(i, j) = \frac{\bar{\mathbf{a}}_i^\top \bar{\mathbf{a}}_j}{\|\bar{\mathbf{a}}_i\| \cdot \|\bar{\mathbf{a}}_j\|} = \frac{\text{pairCount}(i, j)}{\sqrt{\text{itemCount}(i)} \cdot \sqrt{\text{itemCount}(j)}}, \quad j \neq i$$

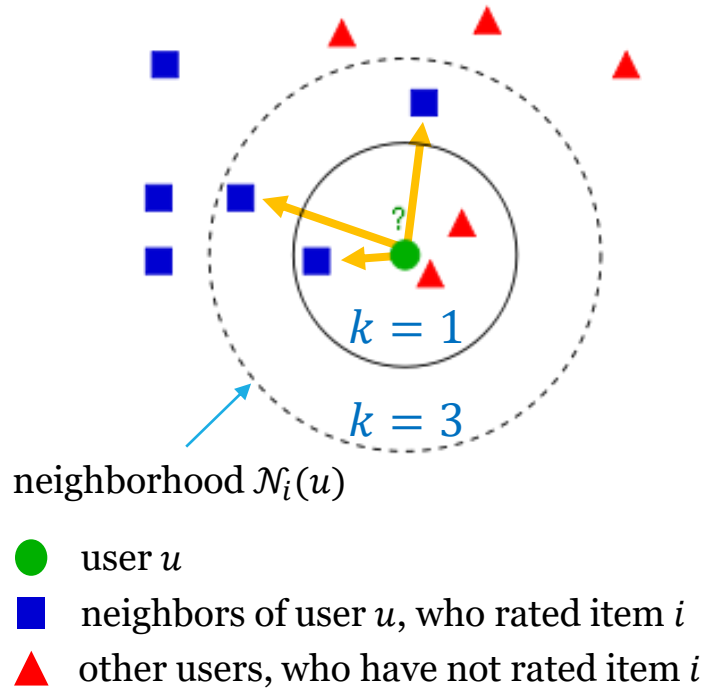
$$\|\bar{\mathbf{a}}_i\|^2 = \sum_u a_{ui}^2 = \sum_u a_{ui} = \text{itemCount}(i) \quad \bar{\mathbf{a}}_i^\top \bar{\mathbf{a}}_j = \sum_u a_{ui} a_{uj} = \text{pairCount}(i, j)$$

After observing ΔA new interactions s.t. $A' = A + \Delta A$, the updated similarity is:

$$\text{sim}'(i, j) = \frac{\text{pairCount}(i, j) + \sum_u [\Delta A]_{ui} [\Delta A]_{uj}}{\sqrt{\text{itemCount}(i) + \sum_u [\Delta A]_{ui}} \cdot \sqrt{\text{itemCount}(j) + \sum_u [\Delta A]_{uj}}}$$

Nearest neighbors models

kNN-based approach



User-based approach

- aggregated opinion of like-minded users:

$$\text{score}_{\text{uKNN}}(u, i) = \text{agg}_{v \in \mathcal{N}_i(u)} a_{vi}$$

Item-based approach:

-

$$\text{score}_{\text{iKNN}}(u, i) =$$

Simple user-based kNN

$$\text{score}_{\text{uKNN}}(u, i) = \frac{1}{|\mathcal{N}_i(u)|} \sum_{v \in \mathcal{N}_i(u)} a_{vj}$$



Potential issues:

- users may have very different interests
- neighborhood size is unlimited

Improved user-based kNN

$$\text{score}_{\text{uKNN}}(u, i) = \frac{1}{z} \sum_{v \in \mathcal{N}_i(u)} \text{sim}(u, v) \cdot a_{vi}$$

$$\mathcal{N}_i(u) = U_i \setminus \{u\}, \quad z = \sum_{v \in \mathcal{N}_i(u)} |\text{sim}(u, v)|$$



Potential issues:

- ~~• other users may have very different interests~~
- large neighborhood size

Dealing with large neighborhood size

- Storing similarities or on-the-fly computations?
- Aggressive subsampling
- Approximate nearest neighbors
 - e.g., NMSLib, Faiss, Annoy
- Dimensionality reduction


Reducing neighborhood

- from N total entities sample $n \ll N$
- select top- k most similar among n samples, $k \ll n$

Possible sampling strategies (must be fast):

- randomly
- most recent only
- most ratings in common (turns into fast MIPS problem)

Example: local time-aware sampling



Sampling strategy:

- select users that have items in common with a target user u
 - each item of a neighbour-user must precede the corresponding item in the target user profile
 - filter out neighbours with too few items in common

Additional weighting:

- users with no recent ratings → lower weights
- active neighbour-user but old rating on a target item → lower weights

	recent user ($t_0 \approx t_{u'1}$)	old user ($t_0 \gg t_{u'1}$)
recent item ($t_{u'1} \approx t_{u'i}$)	≈ 0	$t_0 - t_{u'1}$
old item ($t_{u'1} \gg t_{u'i}$)	$t_{u'1} - t_{u'i}$	$t_0 - t_{u'1}$

Centered kNN

- baseline estimators contain most of useful signal
- every user may have individual “rating scale”

$$\text{score}_{\text{uKNN}}(u, i) = \bar{a}_u + \frac{1}{Z} \sum_{v \in \mathcal{N}_i(u)} \text{sim}(u, v) \cdot (a_{vi} - \bar{a}_v)$$

\bar{a}_u - average rating of user u

item-based kNN:

$$\text{score}_{\text{iKNN}}(u, i) =$$

Centered kNN

- baseline estimators contain most of useful signal
- every user may have individual “rating scale”

user-based kNN:

$$\text{score}_{\text{uKNN}}(u, i) = \bar{a}_u + \frac{1}{Z} \sum_{v \in \mathcal{N}_i(u)} \text{sim}(u, v) \cdot (a_{vi} - \bar{a}_v)$$

\bar{a}_u - average rating of user u

item-based kNN:

$$\text{score}_{\text{iKNN}}(u, i) = \bar{a}_i + \frac{1}{Z} \sum_{j \in \mathcal{N}_u(i)} \text{sim}(i, j) \cdot (a_{uj} - \bar{a}_j)$$

\bar{a}_i - average rating of user i

Similarity measures

- Cosine Similarity
- Pearson Correlation
- Adjusted Cosine Similarity
- Jaccard Index
- Weighted Jaccard Index
- Asymmetric Similarities
- ...
- Spearman's Rank Correlation
- Kendall Tau



insensitive to ranking of “bad” items vs “good ” items

Baseline-adjusted similarity

- **Pearson correlation** (adopted for CF):

$$\text{score}_{\text{Pearson}}(u, v) = \frac{\sum_{i \in I_u \cap I_v} (a_{ui} - \bar{a}_u)(a_{vi} - \bar{a}_v)}{\sqrt{\sum_{i \in I_u \cap I_v} (a_{ui} - \bar{a}_u)^2} \cdot \sqrt{\sum_{i \in I_u \cap I_v} (a_{vi} - \bar{a}_v)^2}}$$

\bar{a}_u - average rating of user u

- **Adjusted Cosine Similarity:**

$$\text{score}_{\text{AC}}(u, v) = \frac{\sum_{i \in I_u \cap I_v} (a_{ui} - \bar{a}_i)(a_{vi} - \bar{a}_i)}{\sqrt{\sum_{i \in I_u \cap I_v} (a_{ui} - \bar{a}_i)^2} \cdot \sqrt{\sum_{i \in I_u \cap I_v} (a_{vi} - \bar{a}_i)^2}}$$


\bar{a}_i - average rating of item i

Jaccard Index

Item-based similarity:

$$\text{sim}_{\text{JI}}(i, j) = \frac{|U_i \cap U_j|}{|U_i \cup U_j|}$$


$$|U_i \cup U_j| = |U_i| + |U_j| - |U_i \cap U_j|$$



Weighted Jaccard Index

- Jaccard Index only operates on sets
- often some values are associated with interactions (e.g., ratings, frequencies)

$$\text{sim}_{\text{WJI}}(u, v) = \frac{\sum_{i=1}^N \min\{w_i(a_u), w_i(a_v)\}}{\sum_{i=1}^N \max\{w_i(a_u), w_i(a_v)\}}, \quad w_i(a_u) = f(a_{ui})$$



	2	1	
	1	3	2

- intersection = []
- union = []
- $\text{score}_{\text{WJI}}(\text{person 1}, \text{person 2}) = \frac{2}{7}$

kNN in matrix form

Element-wise weighting for user-based KNN:

$$r_{ui} = \frac{1}{Z_{ui}} \cdot \sum_{v \in \mathcal{N}_i(u)} \text{sim}(u, v) \cdot a_{vi}, \quad \mathcal{N}_i(u) = U_i \setminus \{u\}$$

we impute A with 0's

kNN in matrix form

Row-wise weighting for user-based KNN:

$$r_{ui} = \frac{1}{Z_{ui}} \cdot \sum_{v \in \mathcal{N}_i(u)} \text{sim}(u, v) \cdot a_{vi}, \quad \mathcal{N}_i(u) = U \setminus \{u\}$$

i.e., “explicit” 0’s

kNN weighting schemes

K – user similarity, $k_{ii} = 0, k_{ij} \geq 0, i \neq j$; S – item similarity matrix, $s_{ii} = 0, s_{ij} \geq 0, i \neq j$.

element-wise weighting:

- User-based:


$$R = KA \oslash (KB)$$

$$b_{ui} = \begin{cases} 1, & \text{if } a_{ui} \text{ is known} \\ 0 & \text{otherwise} \end{cases}$$

- Item-based:

$$R = AS^T \oslash (BS^T)$$

- filters known ratings only
- better for rating prediction



row-wise / no weighting:

- User-based:

$$R = D_K^{-1}KA$$

$$D_K = \text{diag}(K\mathbf{e}) \text{ or } D_K = I$$

$$\mathbf{e} = [1, 1 \dots, 1]^T$$

- Item-based

$$R = AS^T D_S^{-1}$$

$$D_S = \text{diag}(S\mathbf{e}) \text{ or } D_S = I$$

- assumes 0-imputation of unknowns
- better for top- n recommendations

Let's implement simple KNN
models

User-based kNN for top-n recommendations

- Row-wise weighting:

$$R = D^{-1}KA$$

$$r_{ui} = w_u \cdot \sum_{v \in \mathcal{N}(u)} \text{sim}(u, v) \cdot a_{vi}$$

- Is it different from the unweighted case?

- Alternative (column-wise) weighting:

$$R = KD^{-1}A$$

$$r_{ui} = \sum_{v \in \mathcal{N}(u)} \text{sim}(u, v) \cdot w_v \cdot a_{vi}$$

kNN with asymmetric similarity

kNN similarity (e.g., item-based):

row-wise weighted symmetric \rightarrow unweighted asymmetric

$$S_{\text{asym}} = D^{-\alpha} S$$

$$R = A S_{\text{asym}}^{\top}$$

Example: cosine similarity, assuming $d_{ii} = \|\bar{\mathbf{a}}_i\|$:

$$\text{sim}(i, j) = [S_{\text{asym}}]_{ij} = \frac{\bar{\mathbf{a}}_i^{\top} \bar{\mathbf{a}}_j}{\|\bar{\mathbf{a}}_i\|^{1+\alpha} \cdot \|\bar{\mathbf{a}}_j\|}$$

For binary data, $\alpha = -1$ gives a simple conditional probability $p(i|j)$

Popularity effect in asymmetric similarity

$$\text{sim}(i, j) = [D^{-\alpha} S]_{ij} = \frac{\bar{\mathbf{a}}_i^\top \bar{\mathbf{a}}_j}{\|\bar{\mathbf{a}}_i\|^{1+\alpha} \cdot \|\bar{\mathbf{a}}_j\|}$$

What do we recommend?

$$\begin{array}{ccccc} \text{popular item} & \xrightarrow{\text{sim}} & \text{unpopular} & \text{vs.} & \text{unpopular} \xrightarrow{\text{sim}} \text{popular item} \\ i & & j & & i \qquad j \end{array}$$

Popularity effect in asymmetric similarity

popular item $\xrightarrow{\text{sim}}$ unpopular vs. unpopular $\xrightarrow{\text{sim}}$ popular item

$$\alpha < 0$$

- Popular products \rightarrow too trivial recommendations.
- Easy to guess but low value for users + low diversity.

$$\alpha > 0$$

- Recommending niche products increases diversity.
- For users with generic tastes may not fit well.

- Observation: popular items are not very descriptive of users interests.
- Suggest a normalization that would improve item-KNN recommendations.

New scheme – emphasizing contribution of specific user tastes:

$$S_{\text{asym}} = SD^{-\beta}, \quad \text{sim}(i, j) = \frac{\bar{\mathbf{a}}_i^\top \bar{\mathbf{a}}_j}{\|\bar{\mathbf{a}}_i\| \cdot \|\bar{\mathbf{a}}_j\|^{1+\beta}}$$