


Recommender Systems

Lecture 6

Previous lecture: latent factors models

A



N items

- known entries
- unknown entries

$$A_{full} = R + E$$

Task: find utility (relevance) function f_R :

$$f_R: \text{Users} \times \text{Items} \rightarrow \text{Relevance score}$$

Components of the model:

- Utility function $R = PQ^T$, $r_{ij} \approx \mathbf{p}_i^T \mathbf{q}_j$
- Optimization objective defined by $\mathcal{L}(A, R) \rightarrow \min$
- Optimization method (algorithm), e.g. SVD

For top- n recommendations, PureSVD is very efficient and is a strong baseline.

Practical considerations for PureSVD

- SVD for CF is:
 - NOT pure matrix completion
 - NOT pure dimensionality reduction
- rating prediction quality is not important
- scalability
 - folding-in / incremental updates
 - randomized SVD / streaming methods / out-of-memory SVD

Today's lecture

- PureSVD and popularity bias
- Weighted Matrix Factorization
- Netflix Prize Competition
 - Funk SVD
- Optimization methods
 - Stochastic gradient descent
 - Alternating minimization

Mitigating popularity bias in PureSVD


Top singular vector and popularity bias

- Recall, for a simple rank-1 approximation, top-singular value is driven by the common row of ratings

$$A = \mathbf{e} \cdot \mathbf{a}^\top + \epsilon B$$

- More generally, top singular triplet is likely to capture signal mostly from popular items and active users

- Idea: remove it ☺



The effect of removing top singular component

Method	NDCG@50	Recall@50	D@50	Time(min.)
(a)SVD($k = 20$)	0.60597	0.40434	1574	34.8
(b)SVD($k = 19$)	0.60168	0.40088	2139	35.4
(c)SVD($k = 1$)	0.42106	0.19704	290	20.8
SVD($k = 100$)	0.59912	0.37539	2368	88
WRMF($k = 20, \lambda = 10^{-3}$)	0.60678	0.40904	1861.6	214

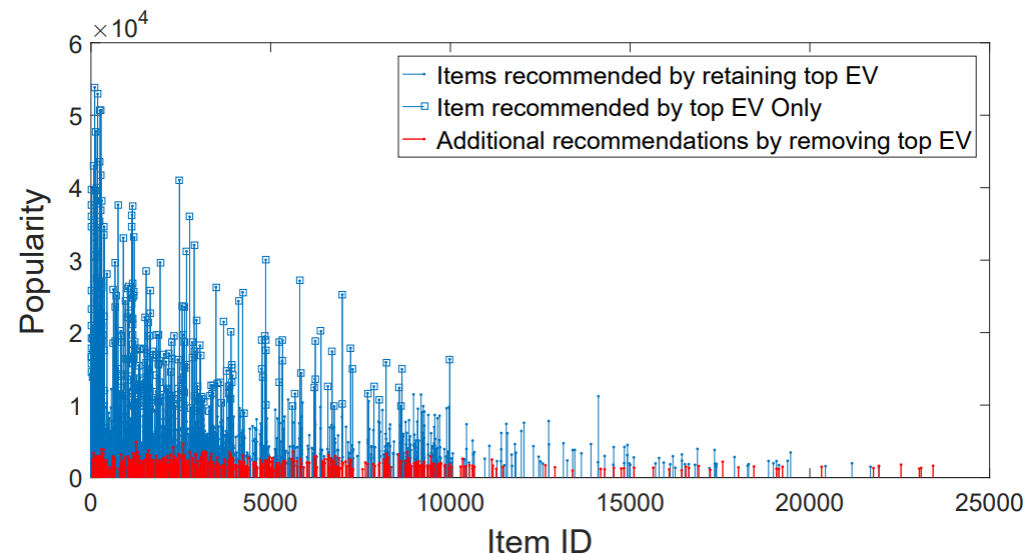


Figure 3: Removing v_H increases diversity by including non-popular items.

Data normalization in PureSVD

- Common observation:
 - interactions data approximately follows power-law or zipf-like distributions
- What effect does it have on covariance matrix?
 - $\mathbf{a}_i^\top \mathbf{a}_j$
- Idea: normalization inversely proportional to popularity

$$\tilde{A} = AD^{f-1}, \quad [D]_{ii} = \|\bar{\mathbf{a}}_i\|$$

Weighted Matrix Factorization

Netflix Prize competition

Contest:

Given a database of movies rated by users,
beat Netflix's recsys **by at least 10%**.

Dates: October 2, 2006 - June 26, 2009

Award: **\$1,000,000**



NETFLIX

Key to success:
ensemble of models.

But latent factors models based
on **matrix factorization**
gained popularity.

Netflix Prize heritage

Impact:

- fueled active research in the field
- signified practical importance for industry
- great organization of the competition

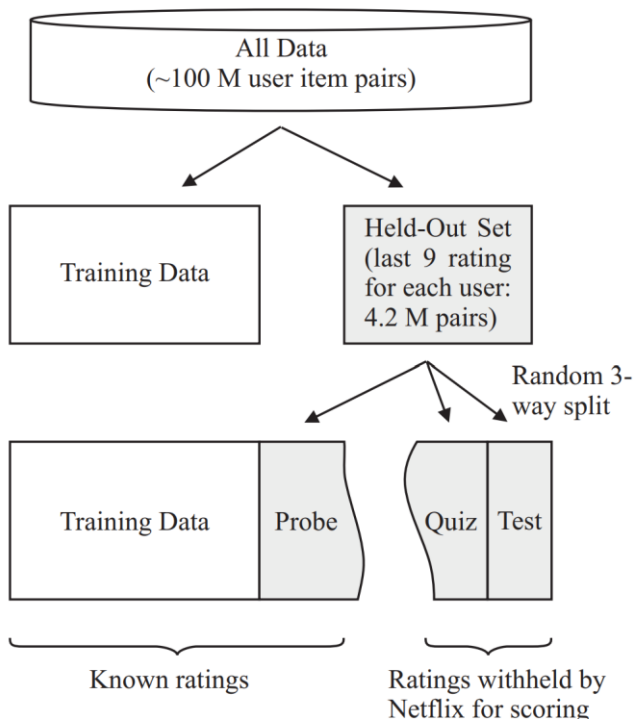


Image source: Takács, Gábor, István Pilászy, Bottyán Németh, and Domonkos Tikk. "Scalable collaborative filtering approaches for large recommender systems." *The Journal of Machine Learning Research* 10 (2009): 623-656.

NETFLIX

Issues:

- treated the problem as a pure matrix completion
- evaluation metric: RMSE
 - high influence on subsequent research echoing to these days
- Actual solution *was never implemented*!
- Users were deanonimized

*<https://www.techdirt.com/blog/innovation/articles/20120409/03412518422/why-netflix-never-implemented-algorithm-that-won-netflix-1-million-challenge.shtml>

Weighted Matrix Factorization

Previously:

- How to handle incomplete data?
- PureSVD: just put 0

Is there a better way?

Weighted Matrix Factorization

Loss function:

$$\mathcal{L}(A, \Theta) = \frac{1}{2} \sum_{i,j \in \mathcal{O}} (a_{ij} - \mathbf{p}_i^\top \mathbf{q}_j)^2$$


$$\mathcal{O} = \{(i, j): a_{ij} \text{ is known}\}$$

Matrix form:

$$\mathcal{L}(A, \Theta) = \|W \odot (A - R)\|_F^2$$

$$R = PQ^\top$$

\odot - Hadamard product



simplest case - binary weights:

$$\begin{cases} w_{ij} = 1, & \text{if } a_{ij} \text{ is known,} \\ w_{ij} = 0, & \text{otherwise.} \end{cases}$$

Recall PureSVD from: $\mathcal{L} = \|A_0 - R\|_F^2$

MF optimization objective

Optimization objective:

$$\mathcal{J}(\Theta) = \mathcal{L}(A, \Theta) + \Omega(\Theta)$$

Model parameters: $\Theta = \{P, Q\}$

$\Omega(\Theta)$ - additional constraints, e.g. L_2 regularization

Typical optimization algorithms:

stochastic gradient descent (SGD)

alternating least squares (ALS)

ALS:

$$\begin{cases} P^* = \arg \min_P \mathcal{J}(\Theta) \\ Q^* = \arg \min_Q \mathcal{J}(\Theta) \end{cases}$$

GD:

$$\begin{cases} \mathbf{p}_i \leftarrow \mathbf{p}_i - \eta \nabla_{\mathbf{p}_i} \mathcal{J} \\ \mathbf{q}_j \leftarrow \mathbf{q}_j - \eta \nabla_{\mathbf{q}_j} \mathcal{J} \end{cases}$$

MF via gradient descent

$$\mathcal{J}(P, Q) = \frac{1}{2} \sum_{i,j \in \mathcal{O}} (a_{ij} - \mathbf{p}_i^\top \mathbf{q}_j)^2 + \lambda (\|\mathbf{p}_i\|^2 + \|\mathbf{q}_j\|^2)$$

Gradient Descent:

$$\begin{cases} \mathbf{p}_i \leftarrow \mathbf{p}_i - \eta \nabla_{\mathbf{p}_i} \mathcal{J} \\ \mathbf{q}_j \leftarrow \mathbf{q}_j - \eta \nabla_{\mathbf{q}_j} \mathcal{J} \end{cases}$$



<https://twitter.com/chaosprime/status/1472385765317521408>

Full gradient:

$$\nabla_{\mathbf{p}_i} \mathcal{J} = - \sum_{j \in I_i} (a_{ij} - \mathbf{p}_i^\top \mathbf{q}_j) \mathbf{q}_j + \lambda \mathbf{p}_i$$

can be inefficient with large data

Optimization with SGD

$$\mathcal{J}(P, Q) = \frac{1}{2} \sum_{i,j \in \mathcal{O}} \overbrace{(a_{ij} - \mathbf{p}_i^\top \mathbf{q}_j)^2}^{l_{ij}} + \lambda (\|\mathbf{p}_i\|^2 + \|\mathbf{q}_j\|^2)$$

Gradient Descent:

$$\begin{cases} \mathbf{p}_i \leftarrow \mathbf{p}_i - \eta \nabla_{\mathbf{p}_i} \mathcal{J} \\ \mathbf{q}_j \leftarrow \mathbf{q}_j - \eta \nabla_{\mathbf{q}_j} \mathcal{J} \end{cases}$$

Idea:

- approximate gradient with its stochastic counterpart
- iterate

Stochastic Gradient Descent:

$$\begin{cases} \mathbf{p}_i \leftarrow \mathbf{p}_i - \eta \frac{\partial l_{ij}}{\partial \mathbf{p}_i} \\ \mathbf{q}_j \leftarrow \mathbf{q}_j - \eta \frac{\partial l_{ij}}{\partial \mathbf{q}_j} \end{cases}$$

$$\frac{\partial l_{ij}}{\partial \mathbf{p}_i} = -(a_{ij} - \mathbf{p}_i^\top \mathbf{q}_j) \mathbf{q}_j + \lambda \mathbf{p}_i$$

Optimization with SGD

$$\frac{\partial l_{ij}}{\partial \mathbf{p}_i} = -(\underbrace{a_{ij} - \mathbf{p}_i^\top \mathbf{q}_j}_{e_{ij}}) \mathbf{q}_j + \lambda \mathbf{p}_i$$

$$\frac{\partial l_{ij}}{\partial \mathbf{q}_i} = -(\underbrace{a_{ij} - \mathbf{p}_i^\top \mathbf{q}_j}_{e_{ij}}) \mathbf{p}_i + \lambda \mathbf{q}_j$$

Algorithm

Initialize P and Q .


Iterate until stopping criteria met:

 for each pair $i, j \in \mathcal{O}$ (shuffled):

 compute e_{ij}

$\mathbf{p}_i \leftarrow \mathbf{p}_i + \eta(e_{ij} \mathbf{q}_j - \lambda \mathbf{p}_i)$

$\mathbf{q}_j \leftarrow \mathbf{q}_j + \eta(e_{ij} \mathbf{p}_i - \lambda \mathbf{q}_j)$




Complexity =

Optimization with SGD

$$\mathcal{J}(P, Q) = \frac{1}{2} \sum_{i,j \in \mathcal{O}} (a_{ij} - \mathbf{p}_i^\top \mathbf{q}_j)^2 + \lambda (\|\mathbf{p}_i\|^2 + \|\mathbf{q}_j\|^2)$$

- What predictions of such model will look like in the case of **binary** a_{ij} ?

Incremental updates



What are the key differences between SGD-based and SVD-based folding-in?

Folding-in

in SVD: $\mathbf{u} = \Sigma^{-1} V^T \mathbf{a}_0$

via SGD:

Initialize \mathbf{p}

Iterate until stopping criteria met:

For all known ratings in \mathbf{a} :

$$e_{aj} = a_j - \mathbf{p}^T \mathbf{q}_j$$

$$\mathbf{p} \leftarrow \mathbf{p} + \eta(e_{aj} \mathbf{q}_j - \lambda \mathbf{p})$$

$$O(\text{nnz}_a \cdot d)$$

of non-zero elements of \mathbf{a}

$$O(\text{nnz}_a \cdot d)$$

Recall - baseline predictors

		
5	3	5
		
3	3	3
		

$$a_{ij} \approx b_{ij} = g_i + f_j + \mu$$

g_i – **g**enerosity of user i , i.e. tendency to assign higher or lower rating

f_j – **f**avoredness of item j , i.e. how likely it's to be praised or critiqued

μ – global average

tends to capture much of the observed signal

Funk SVD: including bias terms into MF



popularized by “Simon Funk”
during the Netflix Prize competition

NETFLIX

- critical users tend to rate movies lower than average user
- popular movies on average receive higher ratings

pre-calculated global average

generosity of a user rating movies higher

$$r_{ij} = \mu + g_i + f_j + \mathbf{p}_i^T \mathbf{q}_j$$

how favorable an item is

$$\underset{\mathbf{p}_i, \mathbf{q}_j, b_i, b_j}{\text{minimize}} \sum_{i,j \in \mathcal{O}} e_{ij}^2 + \lambda (\|\mathbf{p}_i\|^2 + \|\mathbf{q}_j\|^2 + g_i^2 + f_j^2)$$

Iterating over:

$$\begin{cases} \mathbf{p}_i \leftarrow \mathbf{p}_i + \eta(e_{ij}\mathbf{q}_j - \lambda\mathbf{p}_i) \\ \mathbf{q}_j \leftarrow \mathbf{q}_j + \eta(e_{ij}\mathbf{p}_i - \lambda\mathbf{q}_j) \\ g_i \leftarrow t_i + \eta(e_{ij} - \lambda g_i) \\ f_j \leftarrow f_j + \eta(e_{ij} - \lambda f_j) \end{cases}$$

$e_{ij} = ?$

Matrix form

How to incorporate bias terms into a matrix form?

$$[P \ e_M \ t][Q \ f \ e_N]^\top = PQ^\top + e_M f^\top + t e_N^\top$$

Matrix Factorization as unsupervised learning

- Simon Funk - first public solution, Funk SVD.
- Interesting connection to the first attempts to model synaptic connections between neurons (D. Hebb, 1949)

Generalized Hebbian Algorithm

- For multiple input – multiple outputs:

$$\nabla w_{ij} = \eta(y_i x_j - y_i \sum_{k \leq i} w_{kj} y_k)$$

- Can be used to find singular components [Sangers 1989]
 - utilizes Gram-Schmidt process for orthogonalization
- Can be further expanded (omitting orthogonalization) into the form used in FunkSVD

About Simon Funk

Real name :

- Brandyn Webb (<https://sifter.org/~brandyn>)

Employment:

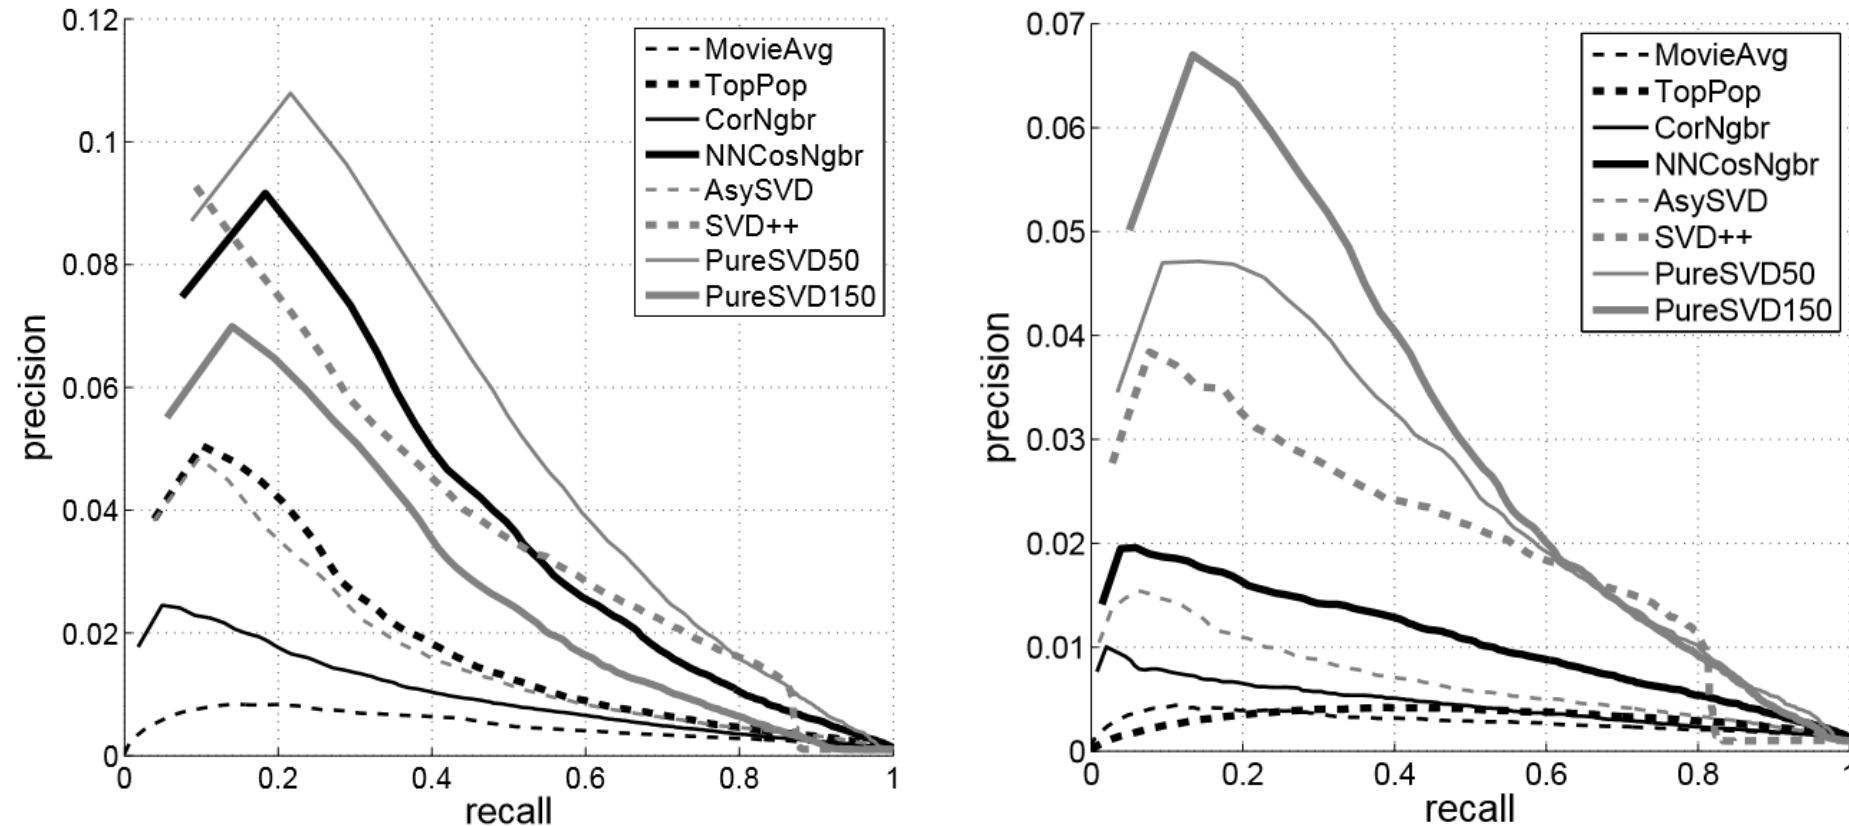
- President of Maui Institute of Cybernetic Epistemology



Source: <https://www.facebook.com/brandyn>

PureSVD vs Weighted MF

Performance on Netflix data: complete dataset (left) and "long-tail" (right).



Images from: P. Cremonesi, Y.Koren, R.Turrin, "Performance of Recommender Algorithms on Top-N Recommendation Tasks", Proceedings of the 4th ACM conference on Recommender systems, 2011.

Note: Funk SVD, SVD++, TimeSVD++, Asymmetric SVD ... **are not** the SVD!

Optimization with Alternating Least Squares

ALS:

$$\begin{cases} P = \arg \min_P \mathcal{J}(\Theta) \\ Q = \arg \min_Q \mathcal{J}(\Theta) \end{cases}$$

$$\mathcal{J}(\Theta) = \mathcal{L}(\Theta) + \Omega(\Theta)$$

$$\mathcal{L}(\Theta) = \frac{1}{2} \|W \odot (A - PQ^\top)\|_F^2$$

$$\Omega(\Theta) = \frac{1}{2} \lambda (\|P\|_F^2 + \|Q\|_F^2)$$

The optimization problem is bi-convex:

“user-oriented” form:

$$\mathcal{J}(\Theta) = \frac{1}{2} \sum_i \|\mathbf{a}_i - Q\mathbf{p}_i\|_{W^{(i)}}^2 + \frac{1}{2} \lambda \sum_i \|\mathbf{p}_i\|_2^2 + \frac{1}{2} \lambda \|Q\|_F^2$$

notation: $\|\mathbf{x}\|_W^2 = \mathbf{x}^\top W \mathbf{x}$

$W^{(i)} = \text{diag}\{w_{i1}, w_{i2}, \dots, w_{iN}\}$

Optimization with ALS

$$\mathcal{J}(\Theta) = \frac{1}{2} \sum_i \|\mathbf{a}_i - Q\mathbf{p}_i\|_{W^{(i)}}^2 + \frac{1}{2} \lambda \sum_i \|\mathbf{p}_i\|_2^2 + \frac{1}{2} \lambda \|Q\|_F^2$$

$$\|\mathbf{x}\|_W^2 = \mathbf{x}^\top W \mathbf{x}$$

$$W^{(i)} = \text{diag}\{w_{i1}, w_{i2}, \dots, w_{iN}\}$$

$$\frac{\partial \mathcal{J}(\Theta)}{\partial P} = 0$$

$$(Q^\top W^{(i)} Q + \lambda I) \mathbf{p}_i = Q^\top W^{(i)} \mathbf{a}_i$$

Optimization with ALS

$$\mathcal{J}(\Theta) = \frac{1}{2} \sum_i \|\mathbf{a}_i - Q\mathbf{p}_i\|_{W^{(i)}}^2 + \frac{1}{2} \lambda \sum_i \|\mathbf{p}_i\|_2^2 + \frac{1}{2} \lambda \|Q\|_F^2$$

$$\|\mathbf{x}\|_W^2 = \mathbf{x}^\top W \mathbf{x}$$


$$W^{(i)} = \text{diag}\{w_{i1}, w_{i2}, \dots, w_{iN}\}$$

$$\bar{W}^{(j)} = \text{diag}\{w_{1j}, w_{2j}, \dots, w_{Mj}\}$$

Block-coordinate descent:

$$\begin{aligned} \frac{\partial \mathcal{J}(\Theta)}{\partial P} &= 0 \\ \frac{\partial \mathcal{J}(\Theta)}{\partial Q} &= 0 \end{aligned}$$

entity-wise updates



$$\begin{aligned} (Q^\top W^{(i)} Q + \lambda I) \mathbf{p}_i &= Q^\top W^{(i)} \mathbf{a}_i \\ (P^\top \bar{W}^{(j)} P + \lambda I) \mathbf{q}_j &= P^\top \bar{W}^{(j)} \bar{\mathbf{a}}_j \end{aligned}$$

system of linear equations:
 $A\mathbf{x} = \mathbf{b}$

“Embarrassingly parallel” algorithm

Initialize P and Q .

Iterate until stopping criteria met:

$$P \leftarrow \arg \min_P \mathcal{J}(\Theta)$$

$$Q \leftarrow \arg \min_Q \mathcal{J}(\Theta)$$

ALS performance

Complexity:

$$O(\quad) + O(\quad)$$

$$(Q^{\top}W^{(i)}Q + \lambda I) \mathbf{p}_i = Q^{\top}W^{(i)}\mathbf{a}_i$$

ALS performance

Complexity:

$$O(\text{nnz}_A \cdot d^2) + O((M + N) \cdot d^3)$$

- can be improved with approximate solvers (e.g. conjugate gradient)
- or switch to coordinate descent method

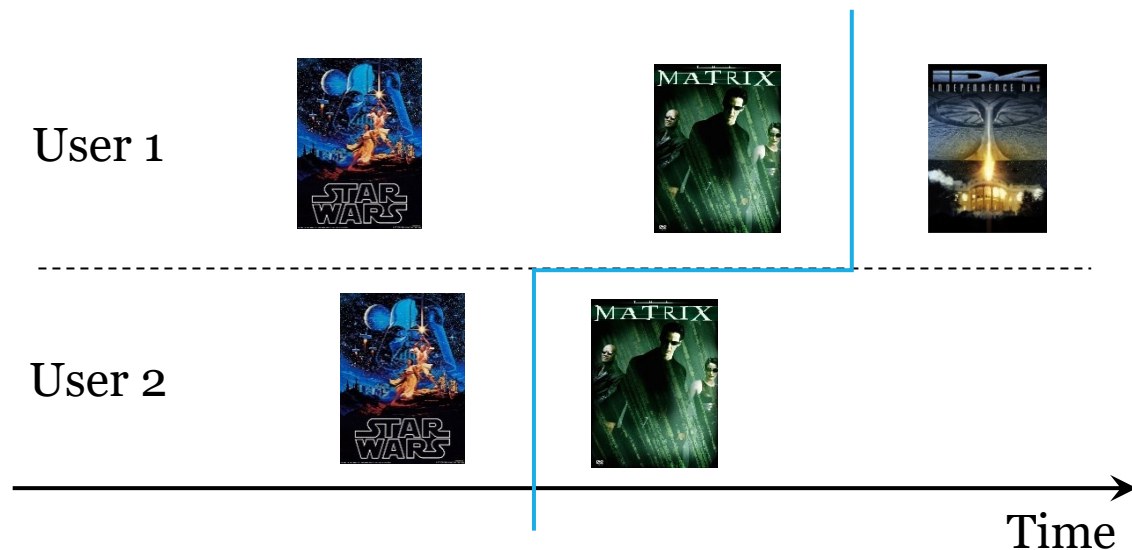
ALS folding-in

Folding-in and timepoint splits

consider scenario:

- warm-start users for evaluation
- test users are the most recently active ones
- most-recent-item holdout

Will there be a data leakage?



Case study: Yandex Zen (old times)

Company manages many different types of media content (news, search, etc.).


Goal: have a unified user representation across all domains.

Solution:

- A neural network embeds onto a latent space all unstructured content
- Users are updated through the “half”-ALS scheme

Algorithm:

- Get Q from external source (ANN)
- Update P based on most recent Q



Case study: Yandex Zen (recent times)

- fully MF-based solution, no neural networks
- solution beats more complex approaches
- SGD-based approach with additional tricks
 - pre-training on larger datasets (prior history)
 - downweight embeddings of “cold” users and items at initialization
 - fp16 calculations

Talk (in Russian)

- Ренессанс факторизации в рекомендациях
<https://www.youtube.com/watch?v=3Inl0iE41NU&list=PLfaEB9oj-8KVBZNFcrESLBGOc5joIxCzK&index=2>

Предобучение

- + Логи за 3 месяца (миллиарды взаимодействий)
- + Hogwild! 5 эпох по часу
- + 60 Gb RAM, 50 ядер

Онлайн дообучение

- + Логи за 5 минут (миллионы взаимодействий)
- + 3 последовательных эпохи за 4 минуты
- + 30Gb RAM, 5 ядер

Practice time

Implement basic MF using SGD for optimization.